Microsoft

# PostgreSQL For SQL Server Professionals

Silvano Coriani
Product @ Azure Postgres

Demos and slides published here:
https://github.com/scoriani/postgres-for-sqlserver-professionals

# Agenda

- Historical notes
- Architectures & Fundamentals
- Deployment options
- Data types & Indexes
- Management
- BCDR
- Programmability
- Security
- App Development
- Cloud Services: Azure Database for PostgreSQL

# SQL Server history

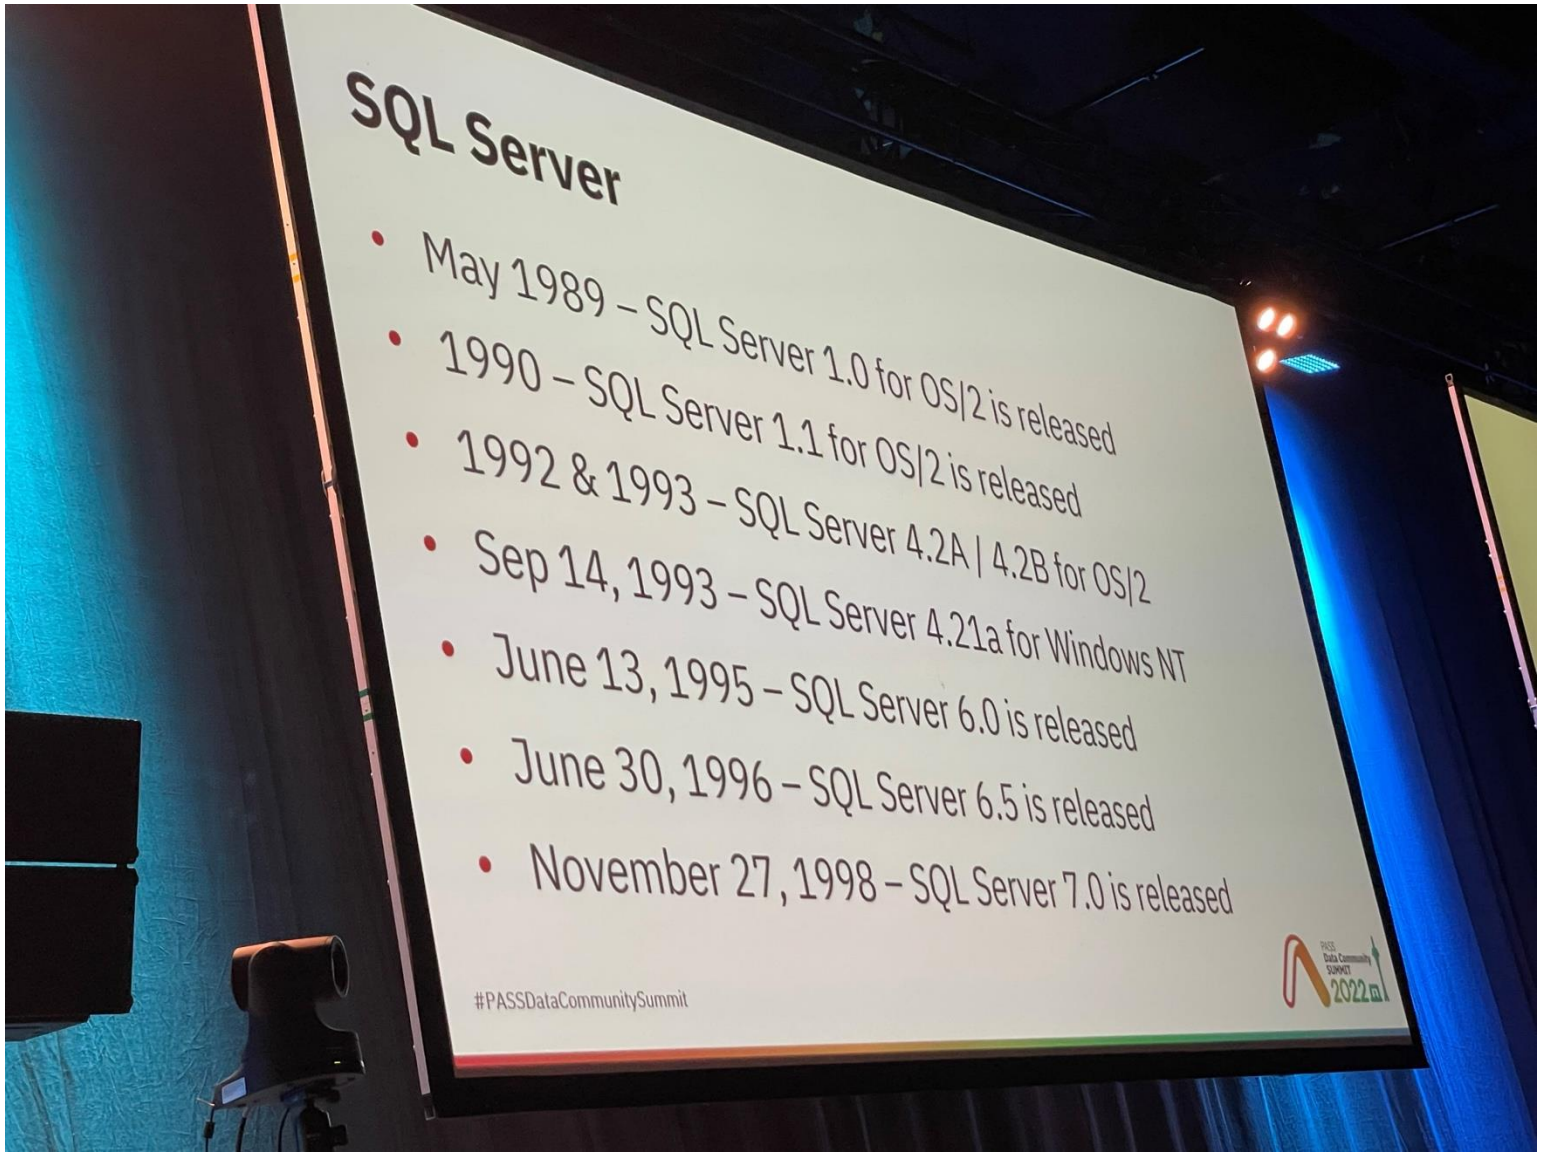

# SQL Server investments over the years

## Continual innovations built into SQL Server, giving you higher return on your investments

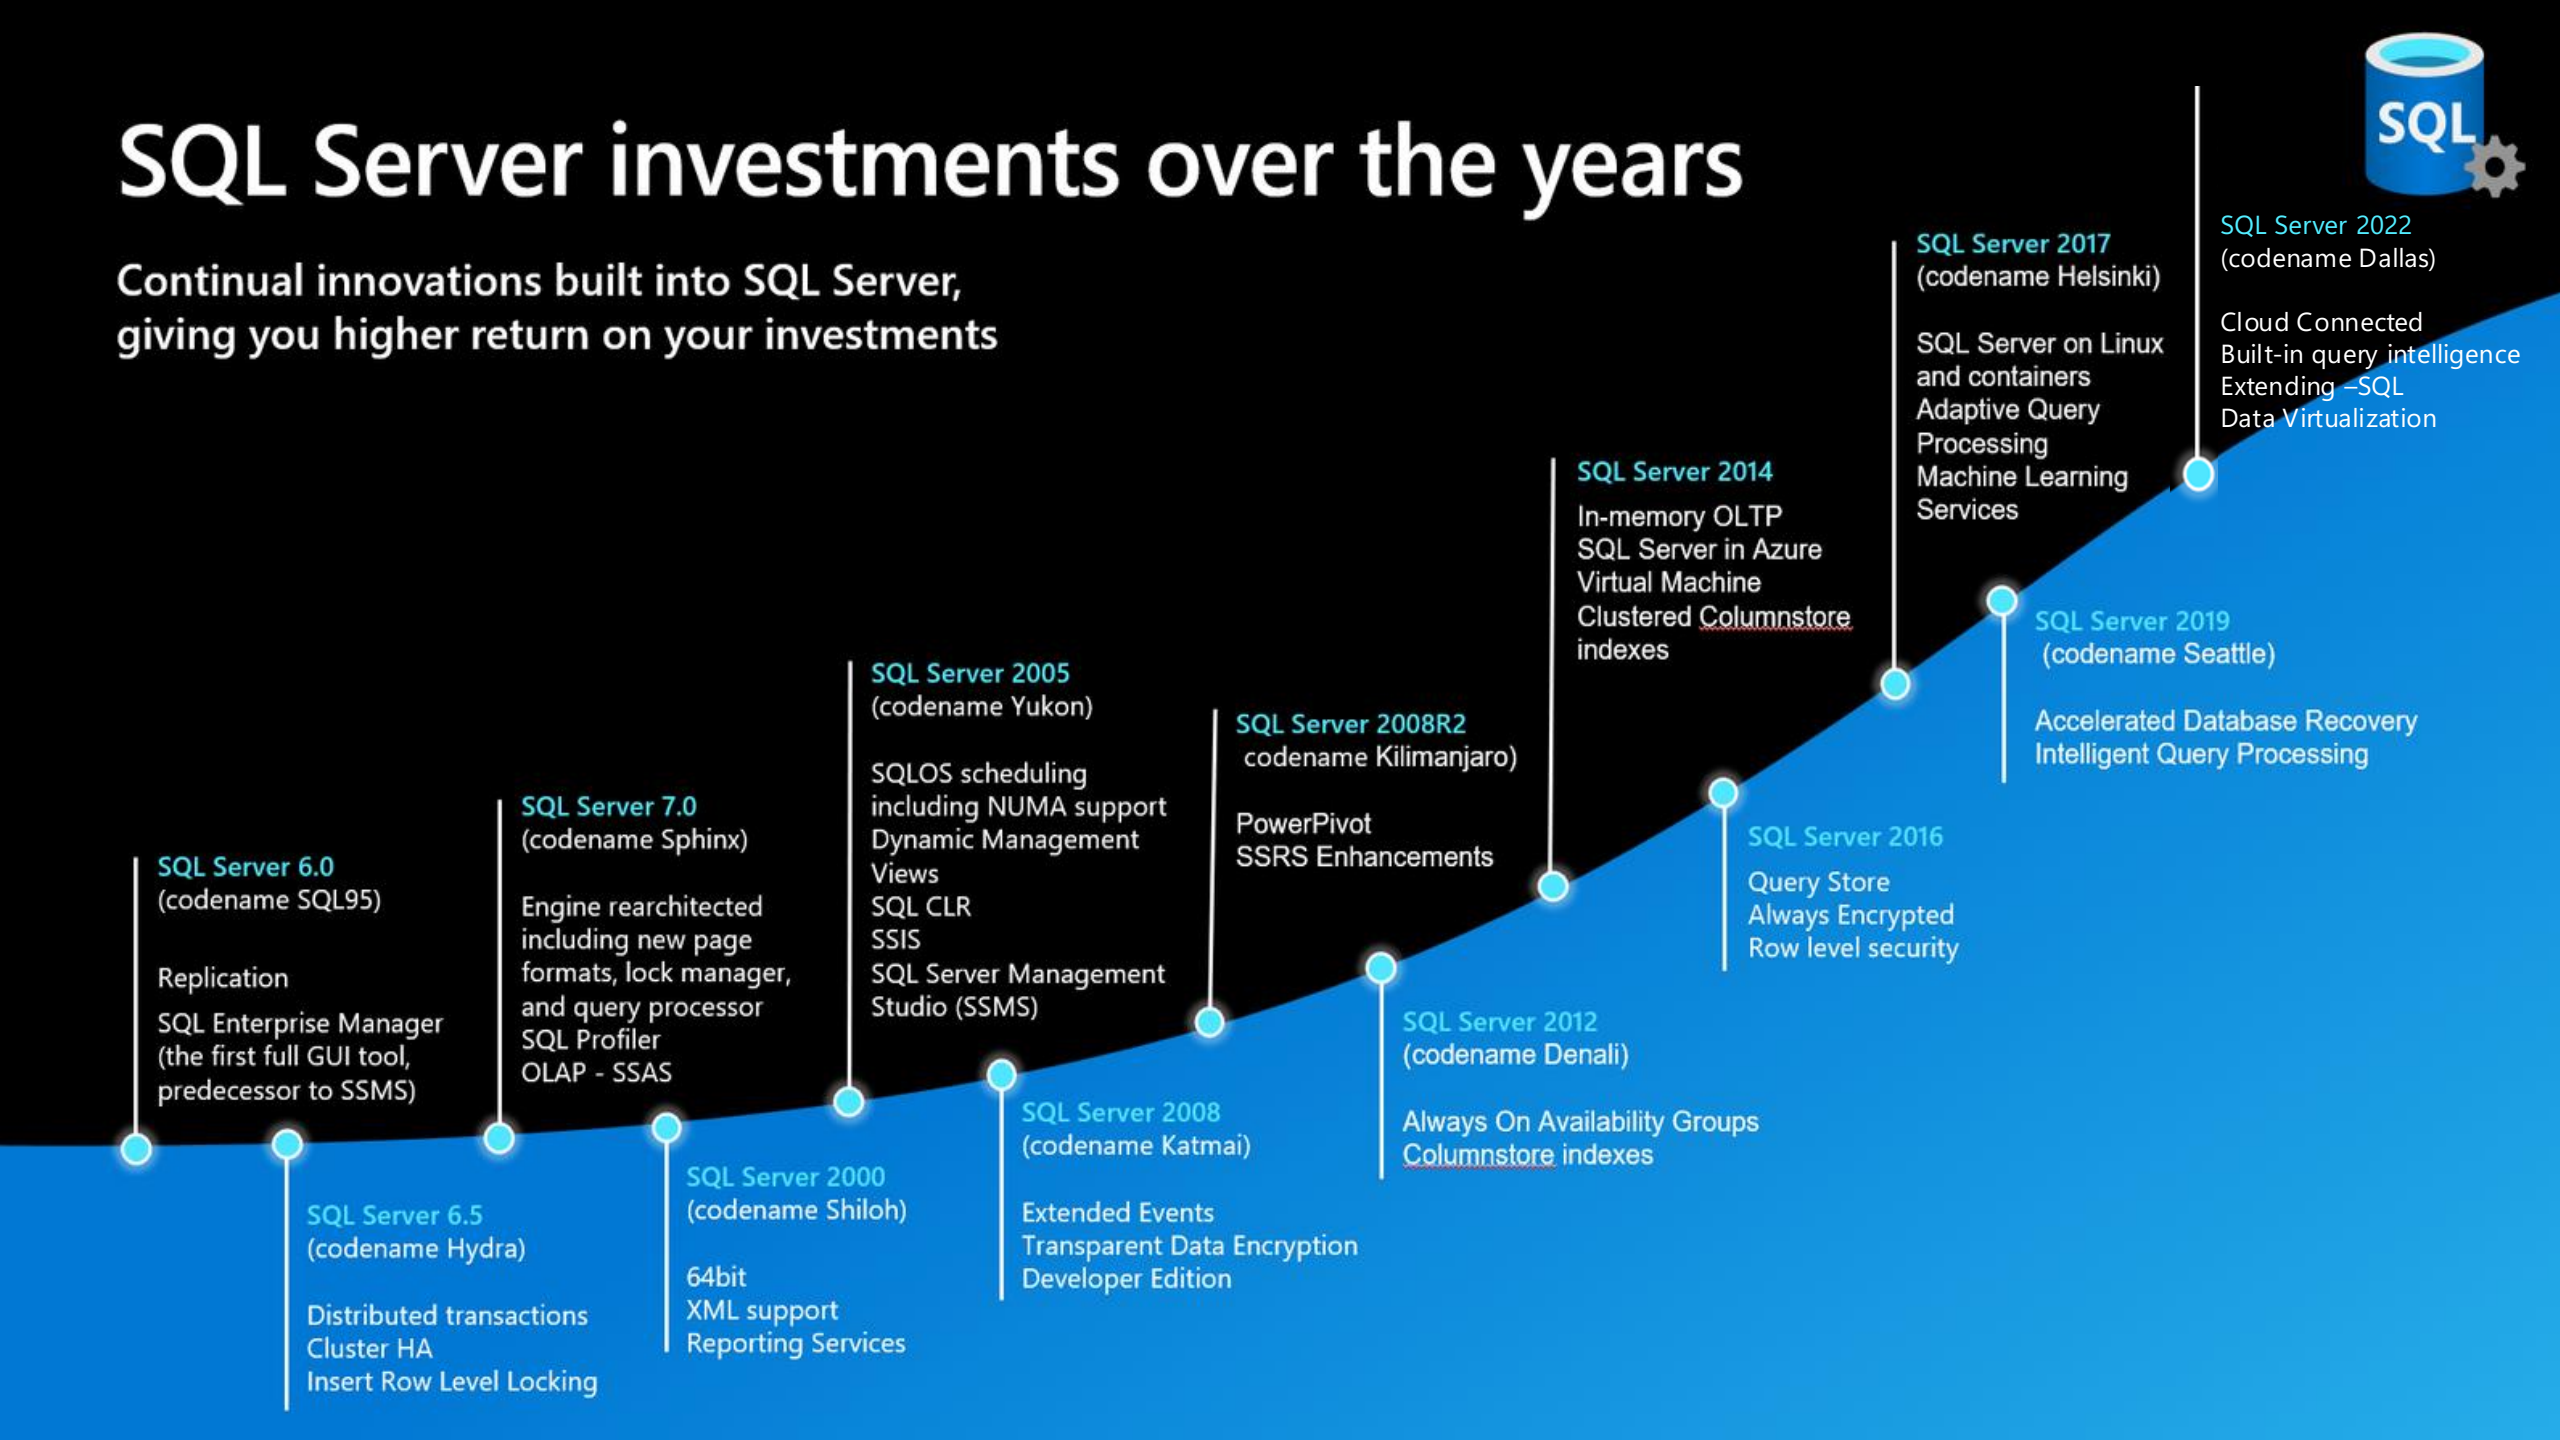

# PostgreSQL History

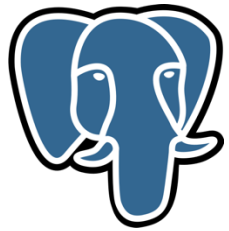


## Inception

- PostgreSQL, initially known as Postgres, evolved from Ingres, a project at the University of California, Berkeley, in the mid-1980s led by Michael Stonebreaker.
- SQL Server (Sybase) and PostgreSQL actually share the same family tree with Ingress being their common ancestor.

## Release

- First released in 1989 as an open-source relational database management system (RDBMS).

## Postgres95

- In 1996, Postgres95 was released, marking a significant milestone with added SQL compliance and enhanced capabilities.

## Community Development

- Community-driven development began in the late 1990s, fostering rapid evolution and feature enhancements.

## Version 7.0

- Released in 2000, introduced several critical features like nested transactions and support for foreign keys.

## Official Name Change

- Renamed as PostgreSQL in 1996 to reflect its support for SQL standards more prominently.

## Ongoing Development

- Continuous development and updates have led to PostgreSQL becoming one of the most advanced and powerful open-source databases available today (v16.2).

# Core engines

# Key fundamentals (1/2)

## SQL Server (proprietary)

- Designed for Windows
- Ported on Linux in 2017 trhough a Platform Abstraction Layer (PAL)
- Supported on x64 proc architectures **for now**
- Multi-threaded
- Built on a concept called SQL Operating System (SQLOS) for
  - user mode thread scheduling
  - memory management
  - synchronization
  - diagnostics/DMVs

## PostgreSQL (open source)

- Community driven OSS , portable
  - Great Community Support and resources
  - Standard compliance
- Designed on Linux/Unix
  - "works" on Windows
  - Multi-process, shared memory
- Runs on all major proc architectures
- Extensibility by design from the core
  - Well-defined APIs to plug-in new implementations (types, operators, etc.)

# SQL Server architecture

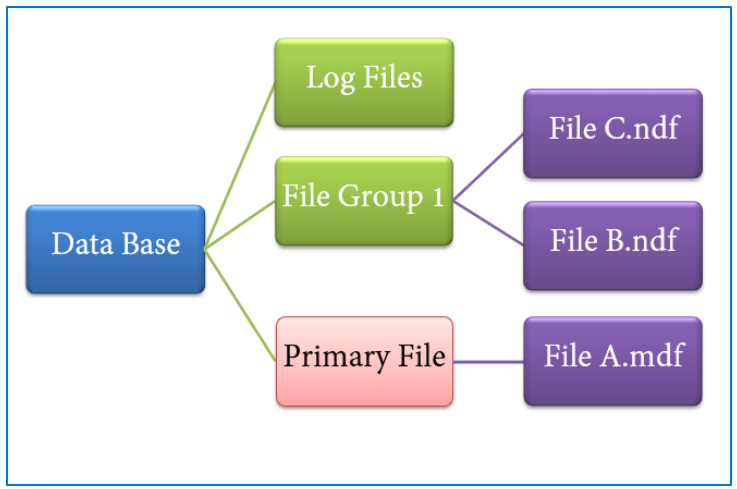


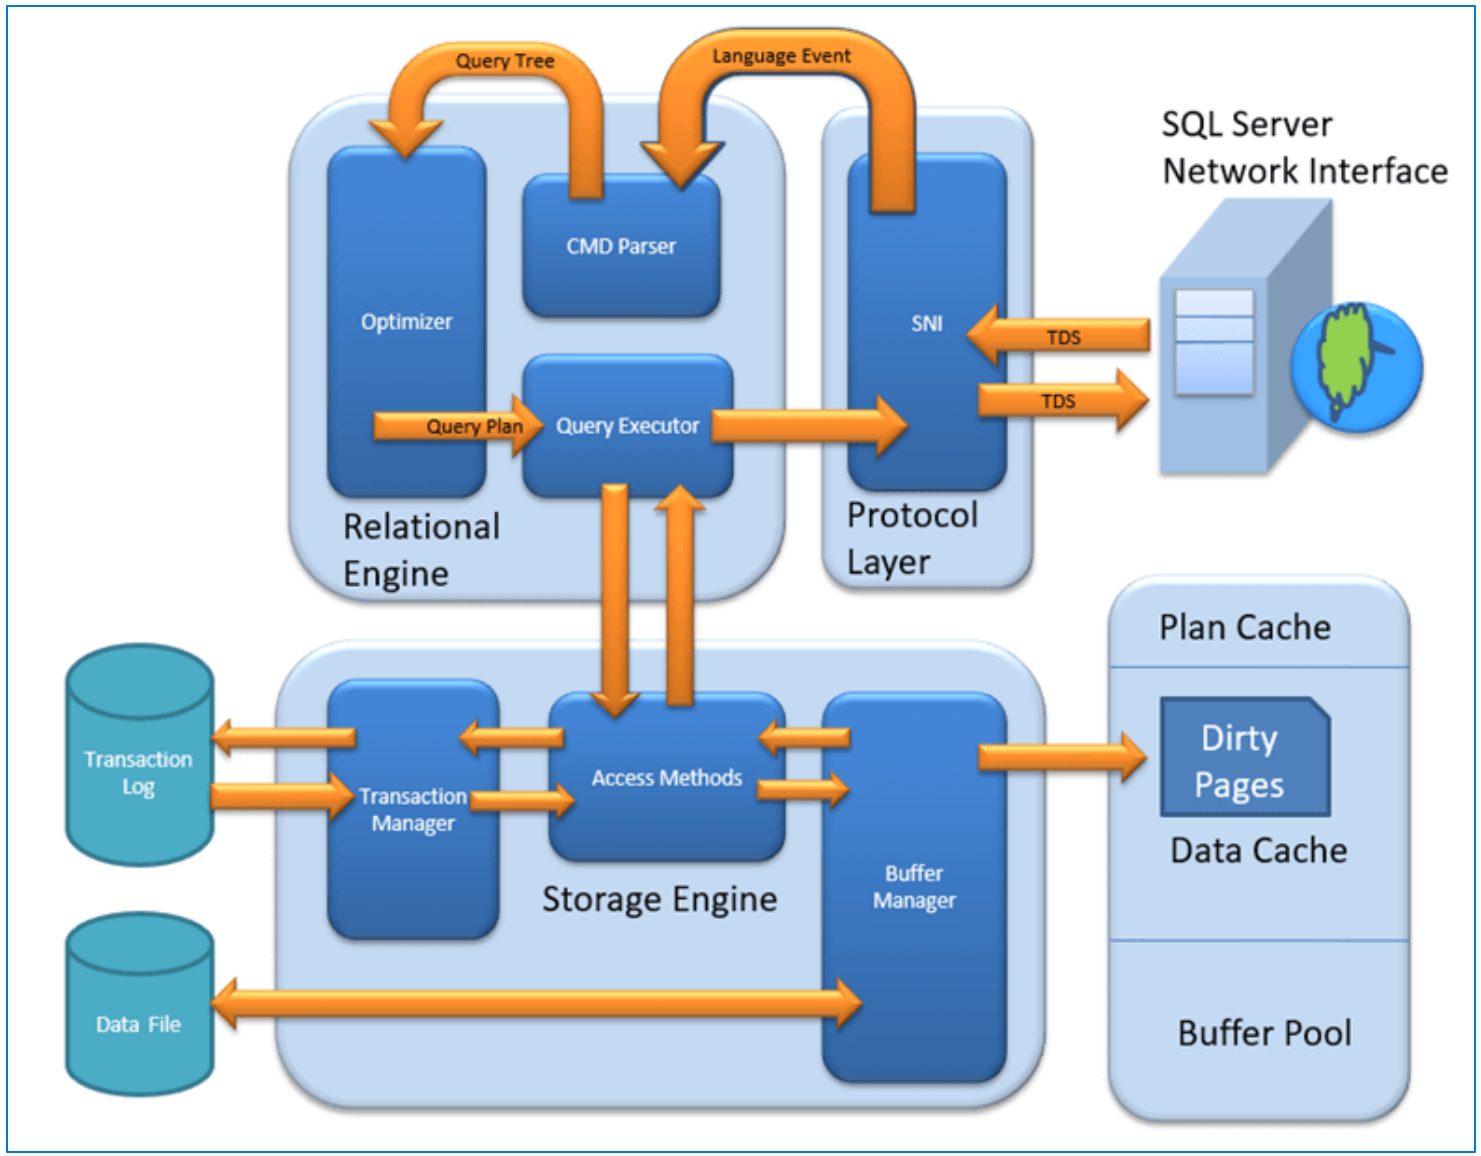


https://www.guru99.com/sql-server-architecture.html

# PostgreSQL architecture

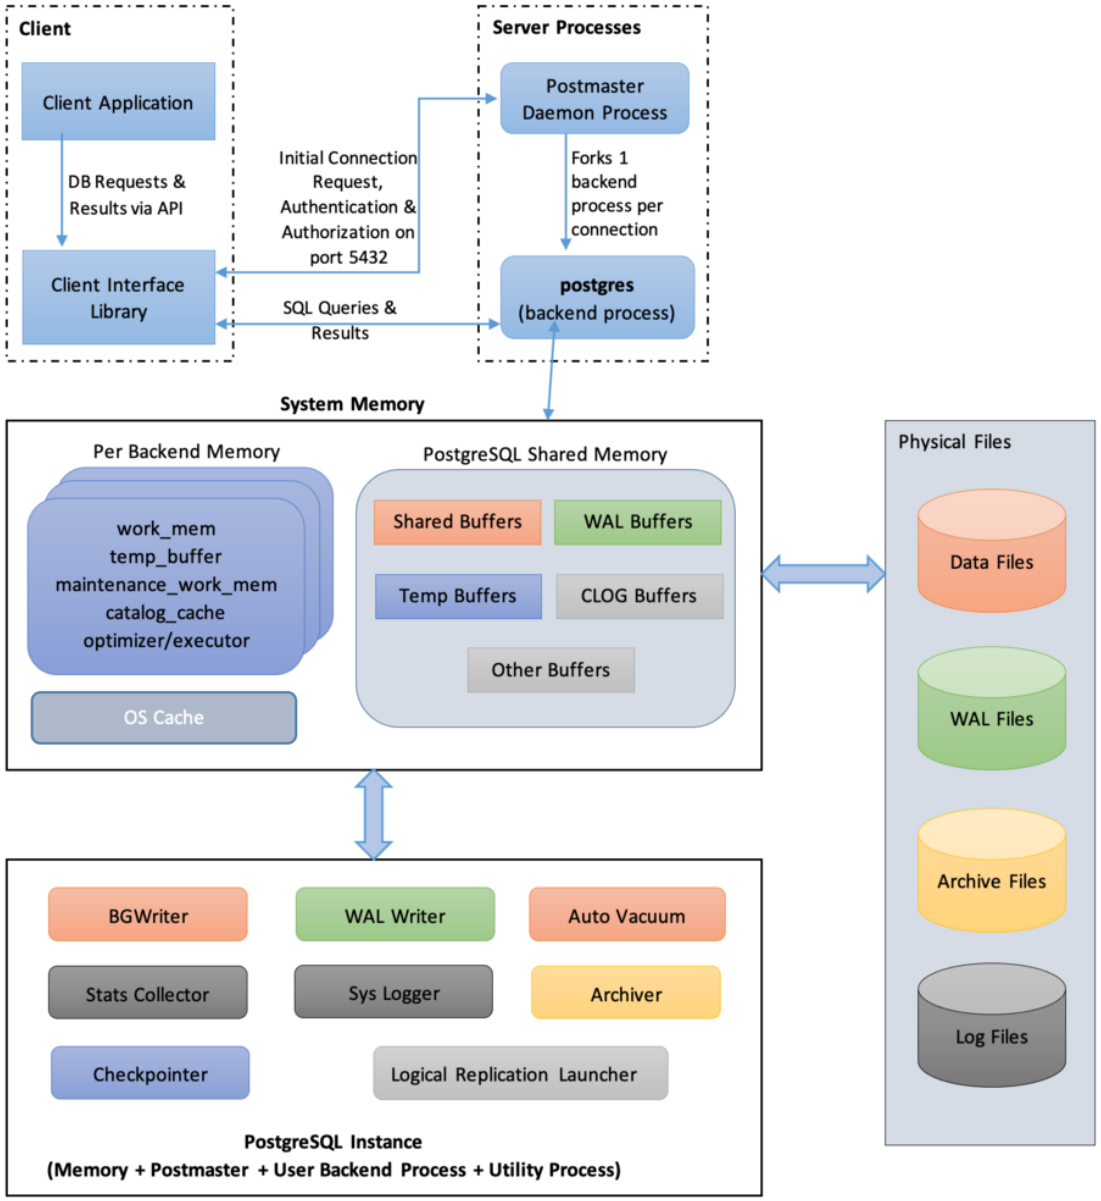

# PostgreSQL process and memory architecture

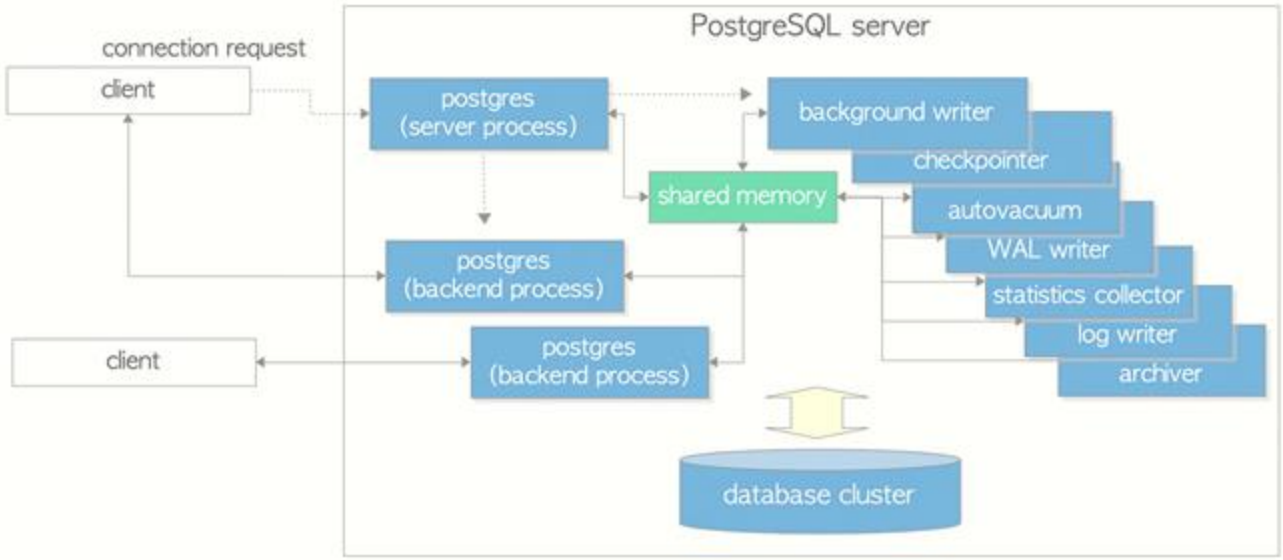


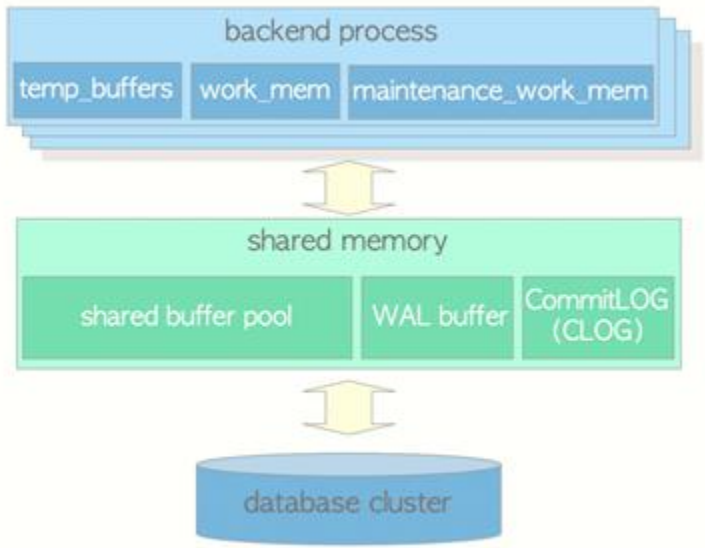


https://www.interdb.jp/pg/pgsql02/01.html

# PostgreSQL storage architecture

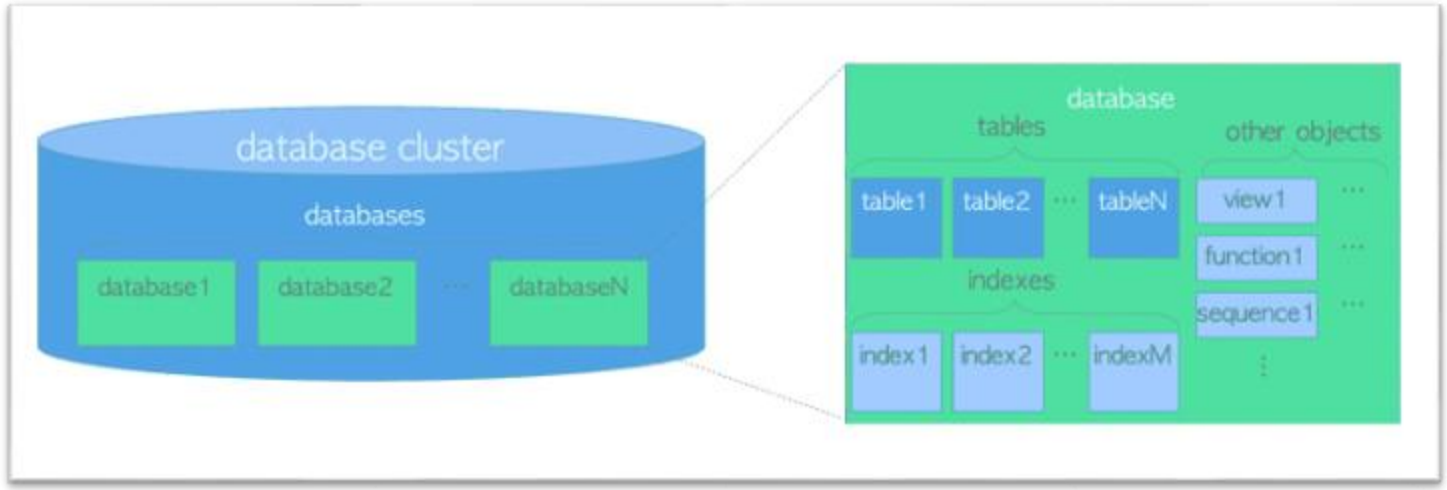


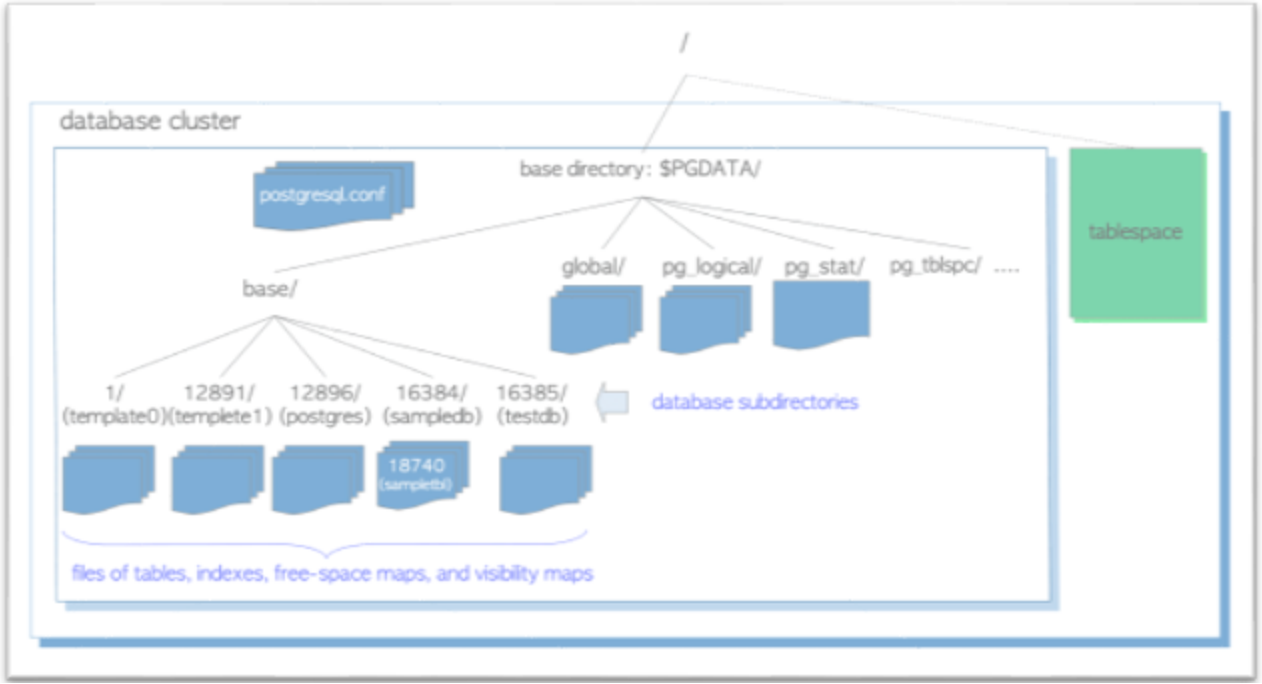


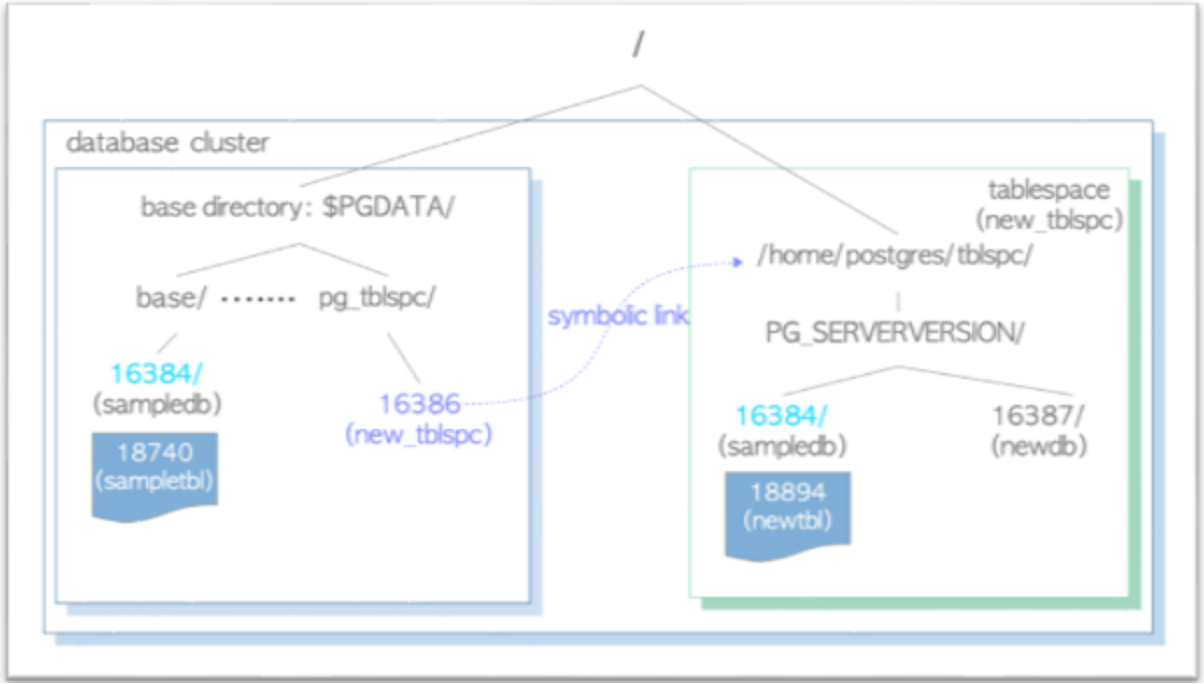

# Key fundamentals (1/2)

## SQL Server (proprietary)

- Designed for Windows
- Ported on Linux in 2017 trhough a Platform Abstraction Layer (PAL)
- Supported on x64 proc architectures only
- Multi-threaded
- Built on a concept called SQL Operating System (SQLOS) for
  - user mode thread scheduling
  - memory management
  - synchronization
  - diagnostics/DMVs

## PostgreSQL (open source)

- Community driven OSS , portable
  - Great Community Support and resources
  - Standard compliance
- Designed on Linux/Unix
  - "works" on Windows
  - Multi-process, shared memory
- Runs on all major proc architectures
- Extensibility by design from the core
  - Well-defined APIs to plug-in new implementations (types, operators, etc.)

# Key fundamentals (2/2)

## SQL Server (proprietary)

- Security / Encryption
- AlwaysON Availability Groups
- AlwaysON Failover Clustering
- Transactional, Merge, P2P replication
- Analysis Services
- Integration Services
- Reporting Services
- SQL Server Management Studio

## PostgreSQL (open source)

- Multi Version Concurrency Control (MVCC)
- Advanced data types (arrays, JSONB, geo, vectors etc.)
- Advanced indexing (GiST, GIN, BRIN, Bloom, etc.)
- Advanced programmability (PL/pgSQL, Perl, Python, Java, etc.)
- High availability with streaming replication
- Logical replication for data integration scenarios

# Deployment options (1/2)

## SQL Server

- Windows Server
  - Native service
  - Unlimited cores and up to 48TB RAM
- Windows 11 dev/test
  - Native service
  - Docker containers
- Linux
  - Supported on RH, SUSE and Ubuntu
  - Containers
  - Kubernetes

## PostgreSQL

- Supported natively on any operating system and proc architecture
  - Local Service or Daemon
  - Containers
  - Kubernetes
- Main distros' package managers
- PostgreSQL.org repositories
- Source code (GitHub or https://git.postgresql.org)
- Config files
  - General service
  - Authentication
  - Identities

# Deployment options (2/2)

## SQL Server

- Cloud Services
  - **Azure SQL**
  - Amazon RDS
  - GCP CloudSQL
  - AliCloud ApisaraDB
  - Others

## PostgreSQL

- Cloud Services
  - **Azure Database for PostgreSQL**
  - **Azure CosmosDB for PostgreSQL**
  - Amazon RDS
  - Amazon Aurora
  - GCP AlloyDB
  - CockroachDB
  - Neon
  - Others

# Demo:

*Deployment options*

# Management

## SQL Server

- Known for ease maintenance efforts
- Maintenance plans
  - Index and stats
  - Backup/restore
  - Consistency checks
- Row versioning (mostly self managing)
  - Snapshot isolation
  - Read Committed Snapshot Isolation (RCSI)
  - Advanced Database Recovery
    - PVS
- SQL Agent

## PostgreSQL

- Routine maintenance tasks
  - Vacuuming
  - Reindexing
  - Backup/restore
  - Log management
- Multi Version Concurrency Control
  - The need for Vacuum process
- Vacuum
  - Concurrent vs full
  - After data load
  - Row changes (including non modification updates)
  - Impact on indexes
  - Autovacuum
  - Autovacuum budget
  - Change thresholds per table
  - HOT UPDATES
  - Inline vacuum
- pgCron

# Monitoring & troubleshooting

## SQL Server

- Activity monitor
- Performance dashboard
- Dynamic Management Views
- Query Store
- Perfmon
- XEvents/Tracing
- Database Tuning Advisor
- Query Tuning Assistant
- Tons of 3rd party tools

## PostgreSQL

- Standard Unix/Linux tools
- Cumulative Statistics System
  - Collection configuration parameters
  - Control functions
  - Predefined views
- Lock contention: pg_locks
- Progress reporting
- pg_stat_statement extension
- Troubleshooting
  - Log analyzer: pgBadger

# Demo:

*Monitoring and troubleshooting*

# BCDR

## SQL Server

- Backup/Restore options
  - Full/Diff/Log
  - Files & Filegroups
  - Compression/Encryption
  - Backup to cloud
- Availability Group topologies
- Failover Clustering
- Log shipping
- Logical Replication
  - Transactional
  - Merge

## PostgreSQL

- Backup/Restore
  - **Logical and physical Backup/restore (pg_dump, pg_basebackup)**
  - File system backup
  - Continuous WAL archiving and Point-in-Time Recovery
- Physical replication for HA (sync/async)
  - Log replication to standby server
  - Failover process
- Logical replication
  - Works for data integration scenarios as well

# Programmability

## SQL Server

- Procedures, functions and triggers
- Spatial data types
- Graph
- .NET CLR
- Language extensibility
- Data virtualization / Linked servers
- Full-text search
- CDC/Change Tracking

## PostgreSQL

- Procedures, functions and triggers
- Schema and Programmability
  - Table inheritance
  - Declarative partitioning
  - Pattern matching ILIKE/SIMILAR TO Regex/POSIX Regex.
- JSON
  - Rich Set of JSON operators
  - Expansive JSON functions
  - Indexing to support pattern matching
- Spatial data types
  - Geometry, geography, raster
  - Indexes
    - Rtree, quadtree
  - Functions
    - ST_Distance, ST_Area, ST_GeometryType,ST_Intersection….
  - Related extensions
- Foreign Data Wrappers
  - SQL Server
  - Oracle
  - …

# Data types

## SQL Server

- Exact numerics
- (Unicode) character strings
- Approximate numerics
- Binary strings
- Date and time
- Spatial
- Hierarchyid
- Rowversion
- UUID

## PostgreSQL

- Numeric
- Monetary
- Character
- Binary Data
- Date/Time
- Boolean
- Enumerated
- Geometric
- Network Address
- Bit String
- Text Search
- UUID
- XML
- JSON
- Arrays
- Composite
- Range
- Domain
- Object Identifier
- pg_lsn
- Pseudo-Types

# Indexes

## SQL Server

- Btree
  - Clustered
  - Non-clustered
- ColumnStore
  - Clustered
  - Non-clustered
- Hash (in memory)
- Filtered
- Computed
- Included columns
- Spatial
- Full-text
- XML

## PostgreSQL

- Indexing & Advanced Indexing
  - Btree
  - Hash
  - GIN
  - GiST, SP-GiST
  - BRIN
  - Partial indexes
  - Index on expressions
  - Extensions can overload access methods to "extend" indexing to other data types (PostgreSQL: Documentation: 16: 38.16. Interfacing Extensions to Indexes)
- pg_trgm - fast searching for similar strings
  - gist_trgm_ops
  - gin_trgm_ops

# Demo:

*Data types and indexes*

# Security

## SQL Server

- Users and Principals
- Ownership and user-schema separation
- Server-level and database-level roles
- Permission hierarchy
- Windows and Azure Active Directory integration
- Encryption
  - Column
  - Storage
  - Always Encrypted
- Auditing
- Ledger

## PostgreSQL

- Robust access-control system
- Client Authentication: GSSAPI, SSPI, LDAP, SCRAM-SHA-256, Certificate, and more
- Database Roles
  - Users
  - Groups
  - Membership
- Column and row-level security
- Encryption
  - Columns
  - Partitions
  - Storage
- Multi-factor authentication with certificates and an additional method
- pgAudit

# Tools

## SQL Server

- Microsoft
  - “new” go-sqlcmd
  - SQL Server Management Studio
  - Azure Data Studio
  - VS Code Extensions
- 3rd party
  - Dev & test
  - Monitoring
  - Management
  - HammerDB
  - …

## PostgreSQL

- Client tools… some!
  - psql
  - pgAdmin
  - Azure Data Studio
  - DBeaver
- PostgreSQL client and server apps
  - reindexdb, vacuumdb, pg_dump, pg_waldump, etc
- Benchmarking
  - pgbench
  - HammerDB
- 3rd party
  - Dev & test
  - Monitoring
  - Management

# App development: drivers and frameworks

## SQL Server

- .NET SqlClient
- Entity Framework Core
- Java
  - JDBC Driver
  - Hibernate
  - Spring Data
- Python
  - pyODBC
  - SQLAlchemy
- Node
  - Tedious
  - Sequelize
- Go
  - go-mssqldb
  - GoORM
- Data API Builder

## PostgreSQL

- libpq — C Library
- libpqxx – C++
- psycopg – Python, SQLAlchemy, Django
- psqlODBC - ODBC
- pjJDBC – Java, Hibernate and Spring
- R2DBC – Reactive Spring
- Npgsql - .NET and Entity Framework Core
- node-postgres – Node.js
- Sequalize and pg
- 3rd party
  - Prisma
  - Devart
  - Debezium
  - Many others
- **Data API Builder!**

# Demo:

*Tools and App development*

# Cloud services

# Azure builds upon PostgreSQL

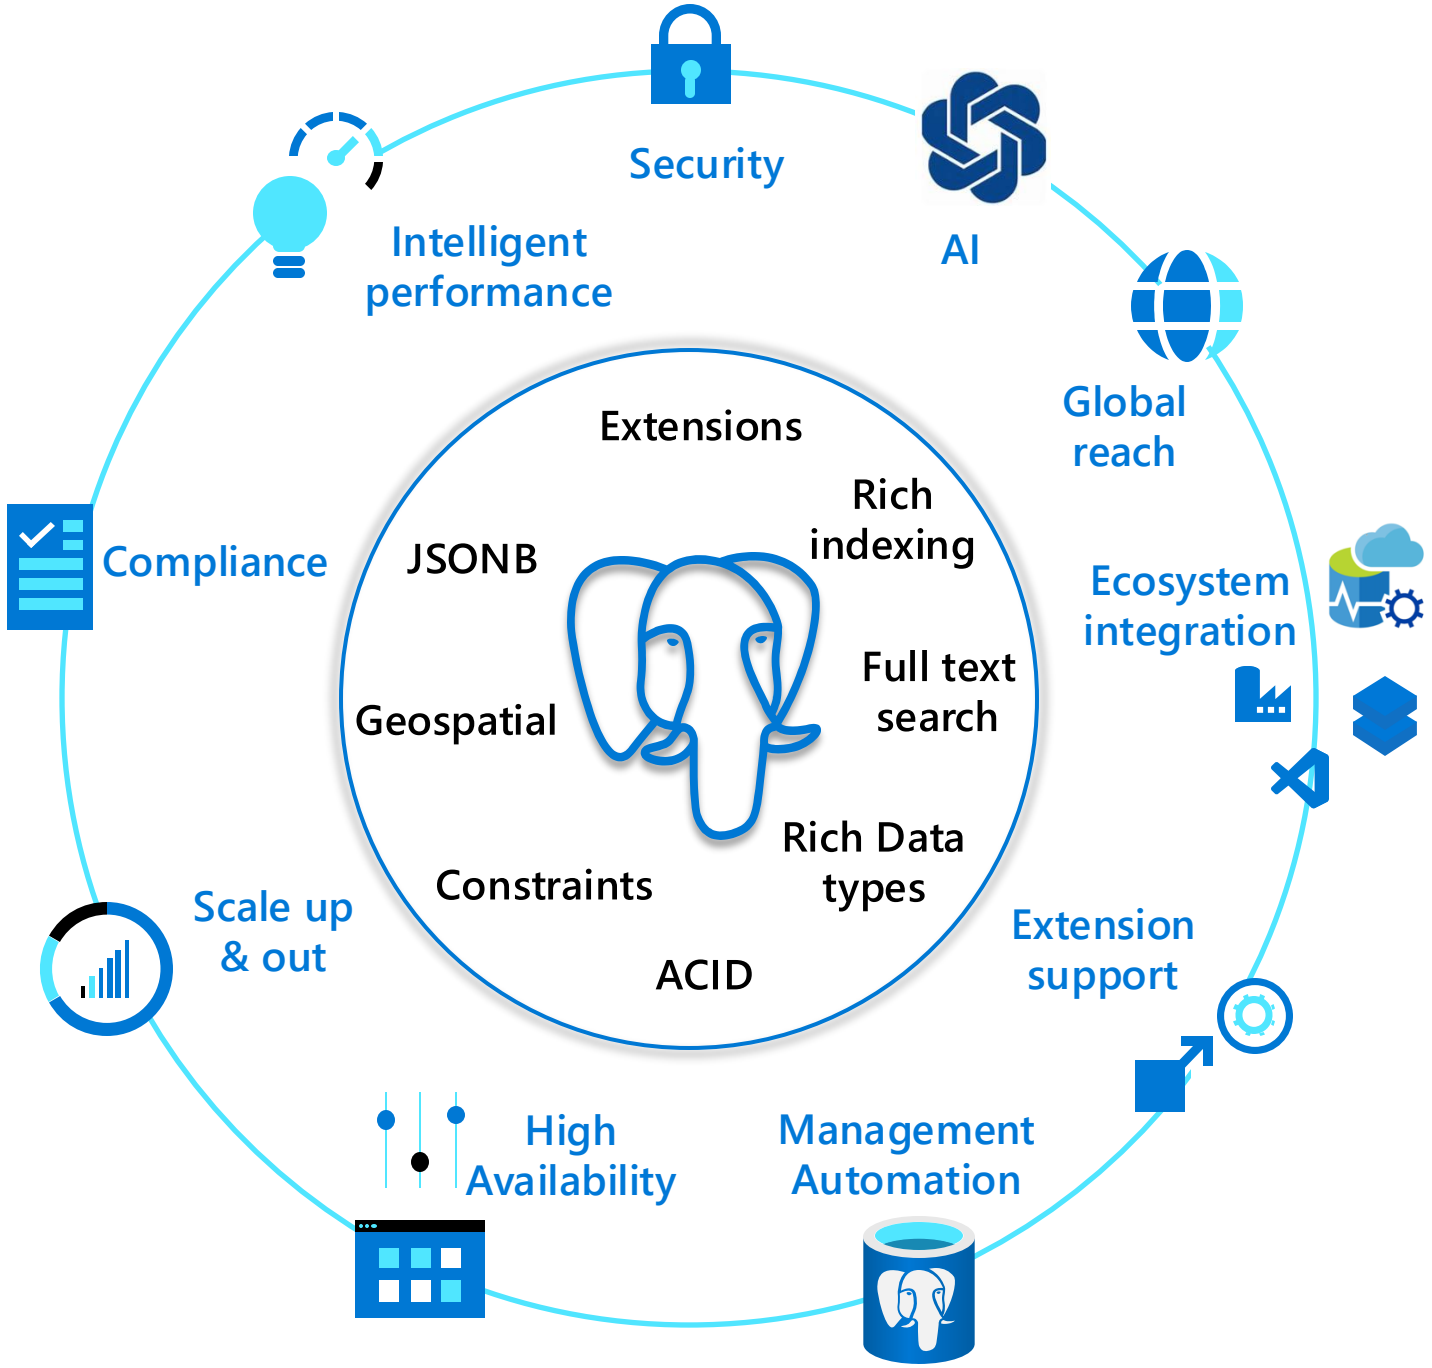

# Flexible Server Architecture

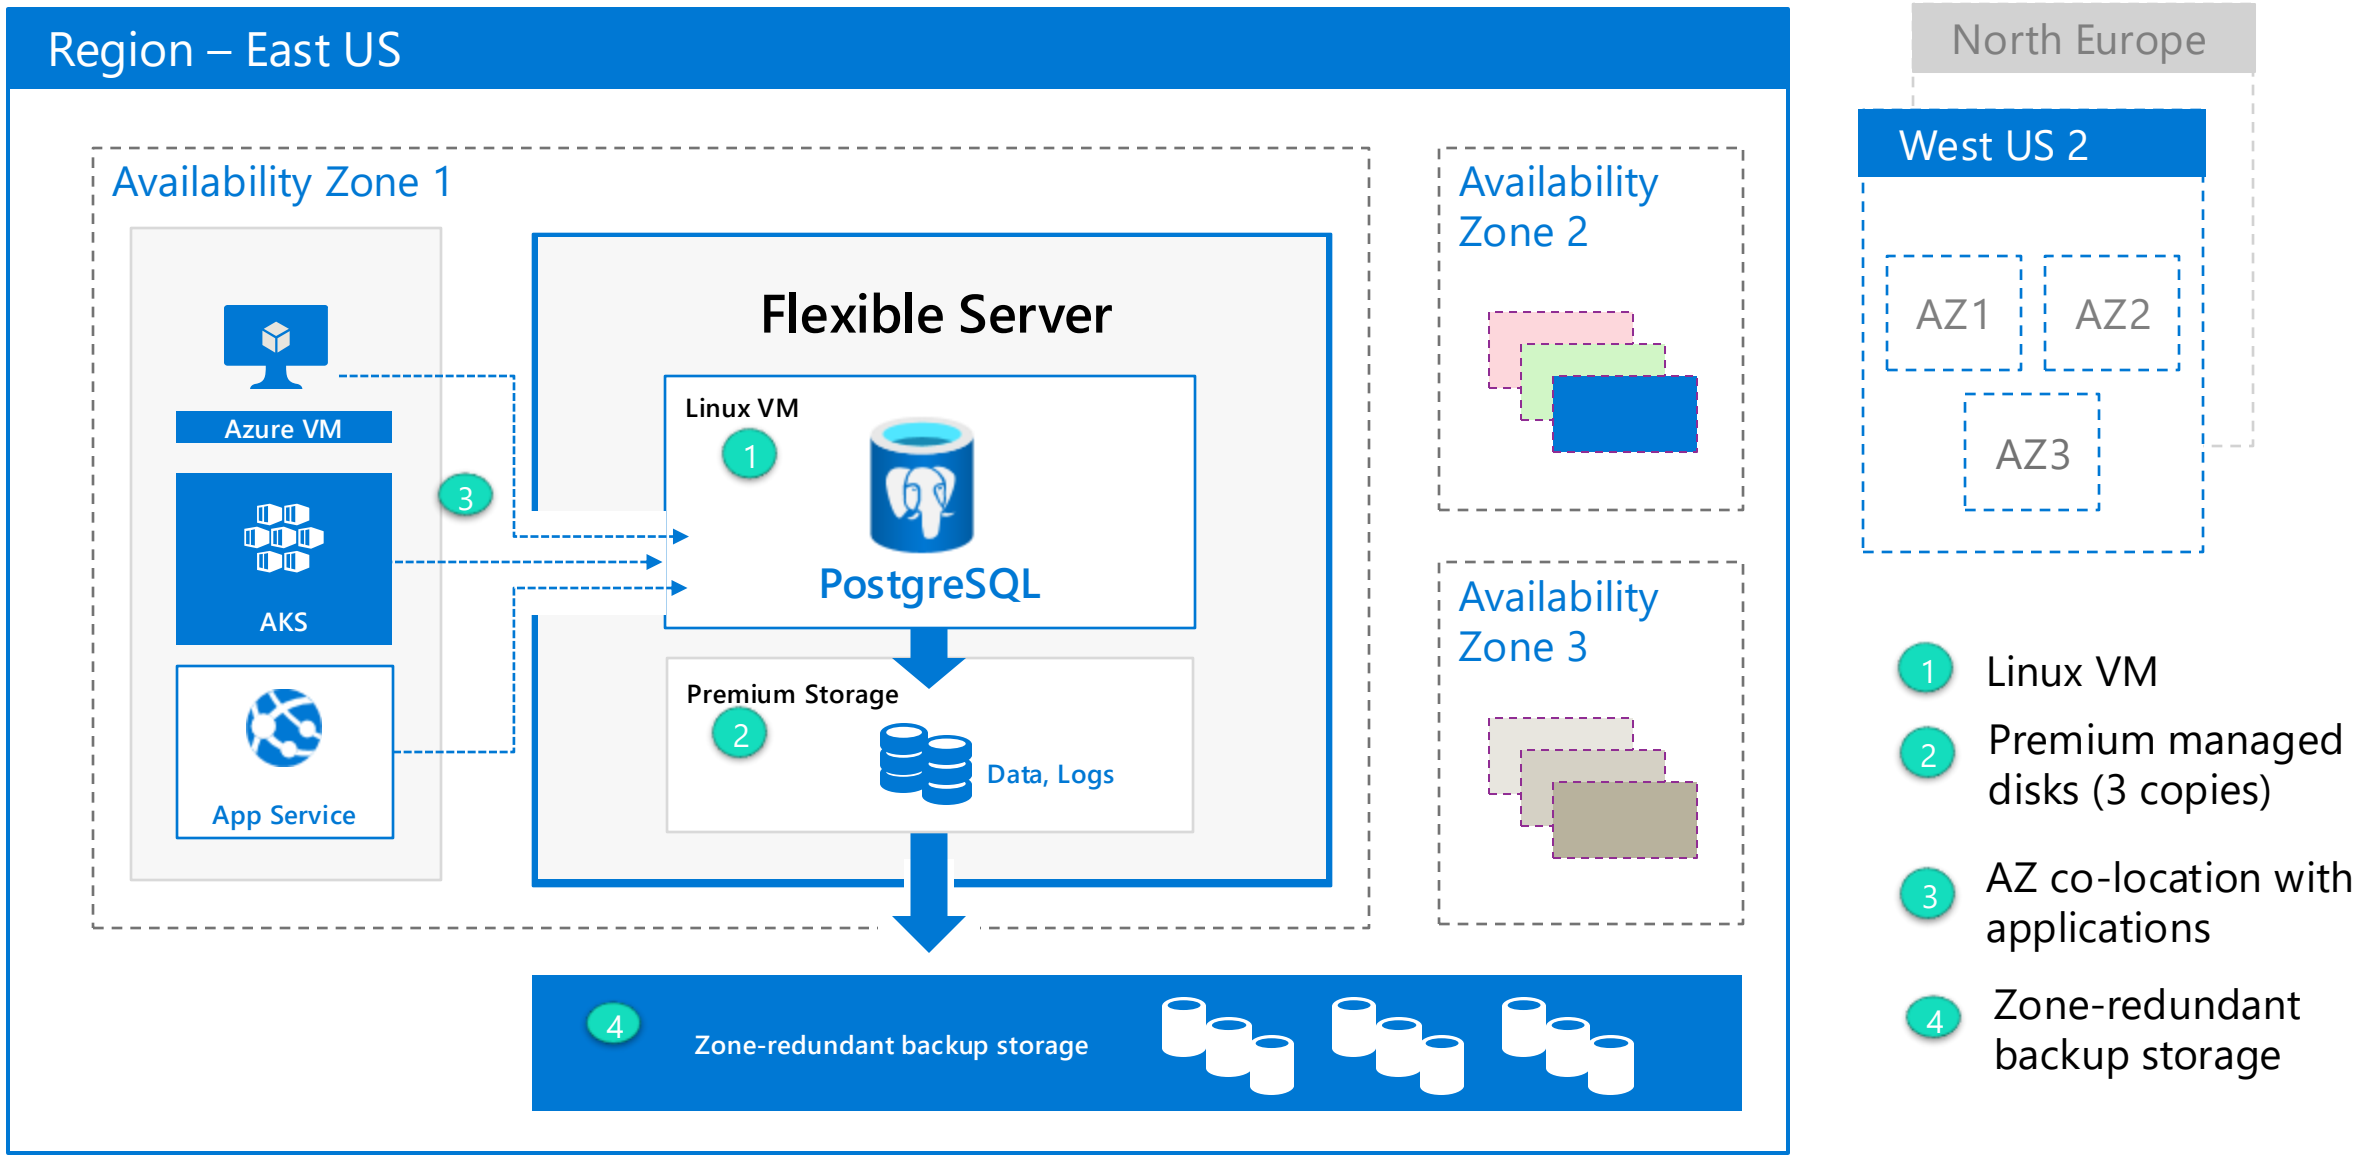

# Workload Optimized Compute SKU's

Cost optimized for different workloads

Each switch between any SKU in minutes

Stop/Start during inactive periods

Reserved Capacity

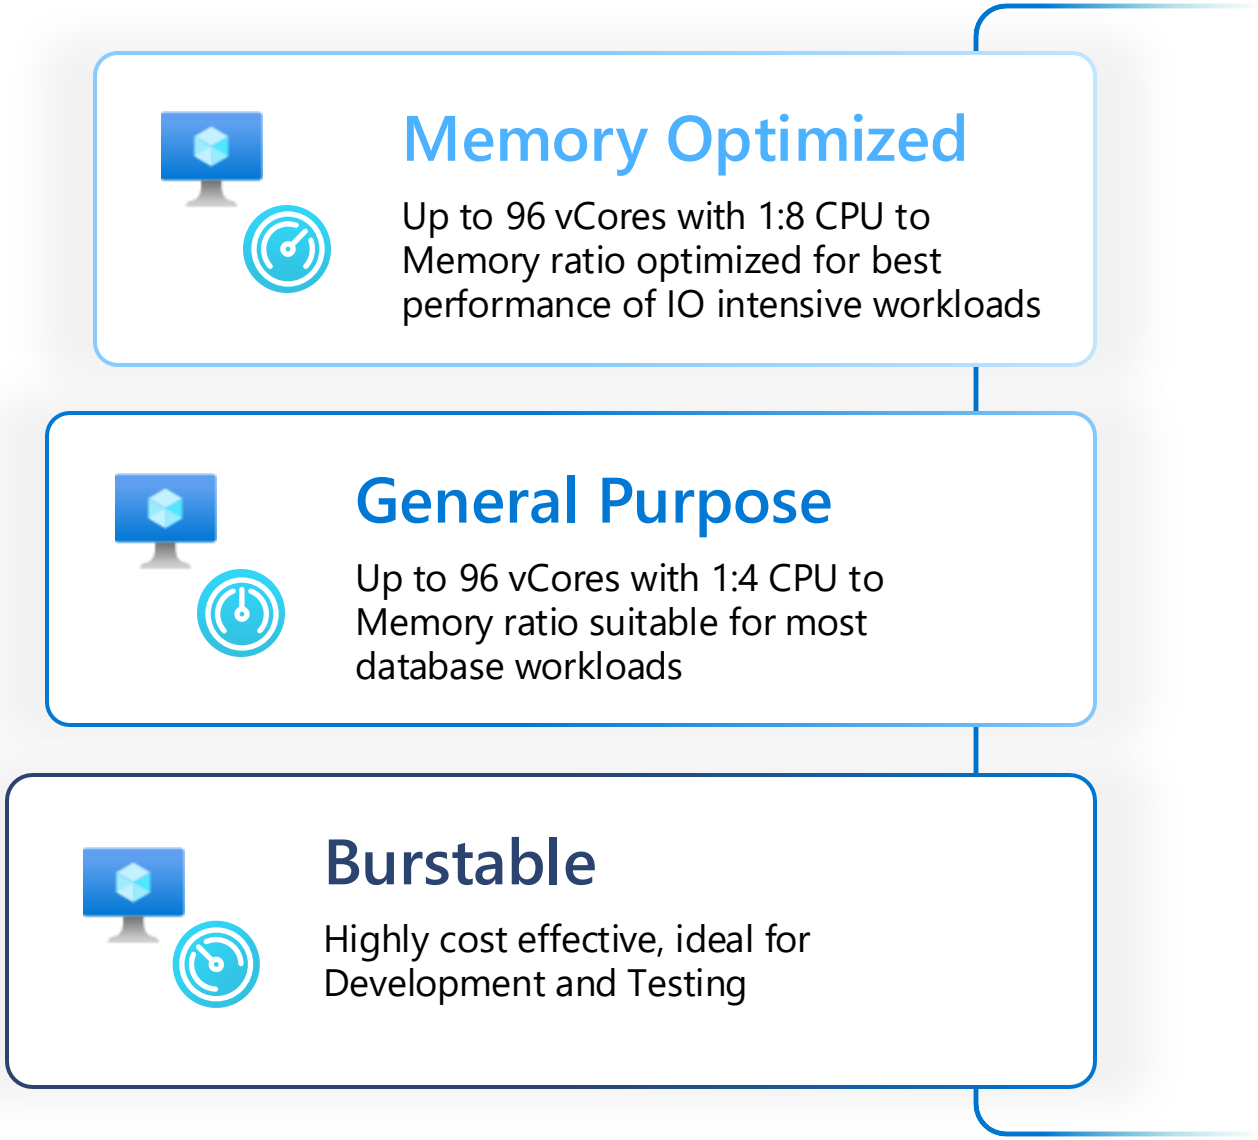

# Elastic Compute and Storage

Scale compute in **under 20s**

Scale storage **online**

Scale **storage size and IOPS** separately

**99.99% SLA** with Availability Zones

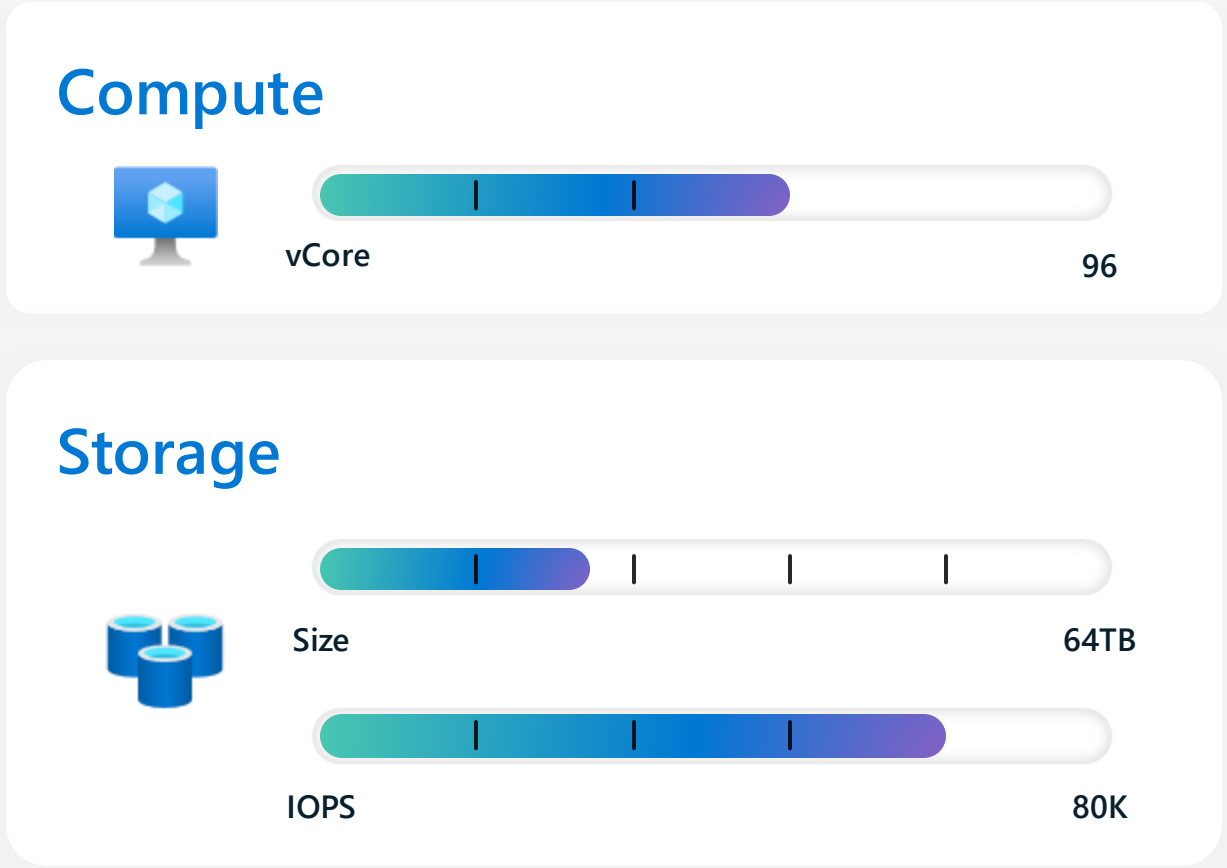

# All Community Supported Postgres Versions

Major versions available **with weeks** of community GA

Minor versions **maintained automatically**

**Upgrade in-place** in minutes

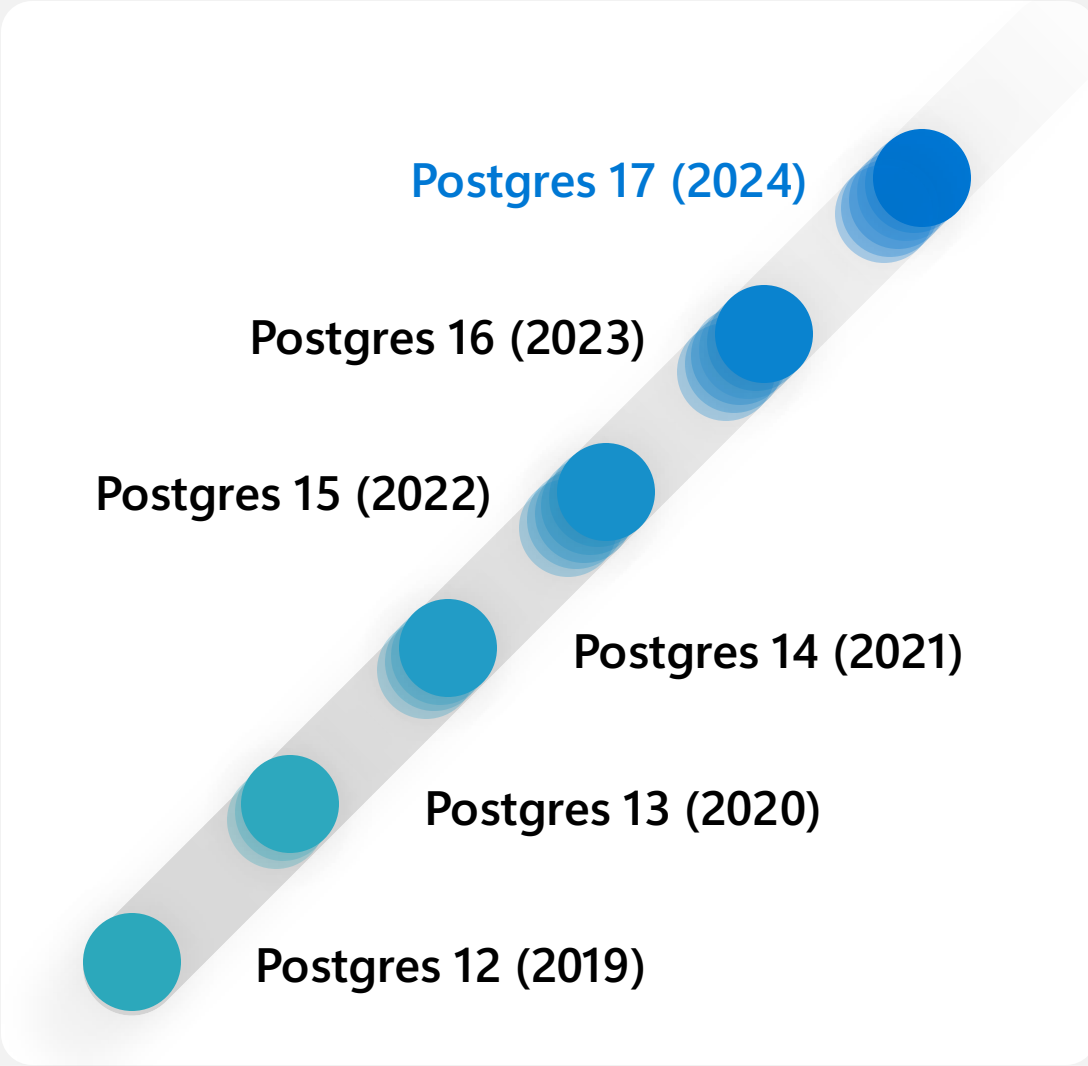

# Broad support for common Postgres extensions

**60+ Postgres extensions** supported

Enables developers to **extend the functionality of Postgres** beyond core capabilities

Microsoft **automatically maintains** extensions versions

address_standardizer
address_standardizer_data_us
amcheck
azure_ai
azure_local_ai (Preview)
**azure_storage**
bloom
btree_gin
btree_gist
citext
cube
dblink
dict_int
dict_xsyn
earthdistance
fuzzystrmatch
hstore
hypopg
intagg
intarray
isn
lo
login_hook
ltree
orafce
pageinspect
**pgaudit**
pg_buffercache
**pg_cron**
pgcrypto
pg_failover_slots (Preview)
pg_freespacemap
pg_hint_plan
pglogical
**pg_partman**
pg_prewarm
pg_repack
pgrouting
pgrowlocks
pg_squeeze
pg_stat_statements
pgstattuple
pg_trgm
pg_visibility
plpgsql
plv8
**postgis**
postgis_raster
postgis_sfcgal
postgis_tiger_geocoder
postgis_topology
postgres_fdw
semver
session_variable
sslinfo
tablefunc
tds_fdw
timescaledb
tsm_system_rows
tsm_system_time
unaccent
uuid-ossp
vector

# Extensive Monitoring

**Rich metrics and logs provide observability into the entire database workload**

Access to detailed metrics and logs

Quickly **diagnose performance** issues

Make informed scaling decisions

**Set up alerts** and auto-scaling for quick responses

**Visualize data** using the Portal, Power BI, Grafana, or Log Analytics

# Enterprise Security

Azure Database for PostgreSQL is **the only Postgres offering** with support for Entra Id Authentication

Service or Customer Managed encryption keys to **protect data at rest**

**Secure network connectivity** via Private Endpoints

## Entra Id

Enterprise Identity built into Postgres

## Customer Managed Keys

Key Vault integration with optional HSM support

## Private Endpoints

Network interface that uses a private IP address from your virtual network

# High Availability

Provides a replica of the production database across Availability Zones with a (no data loss) **Recovery Point Objective = 0**

**Automatic Failover** in scenarios where servers or zones fail

**Automatic rebuild** of the Secondary HA instance after failover occurs

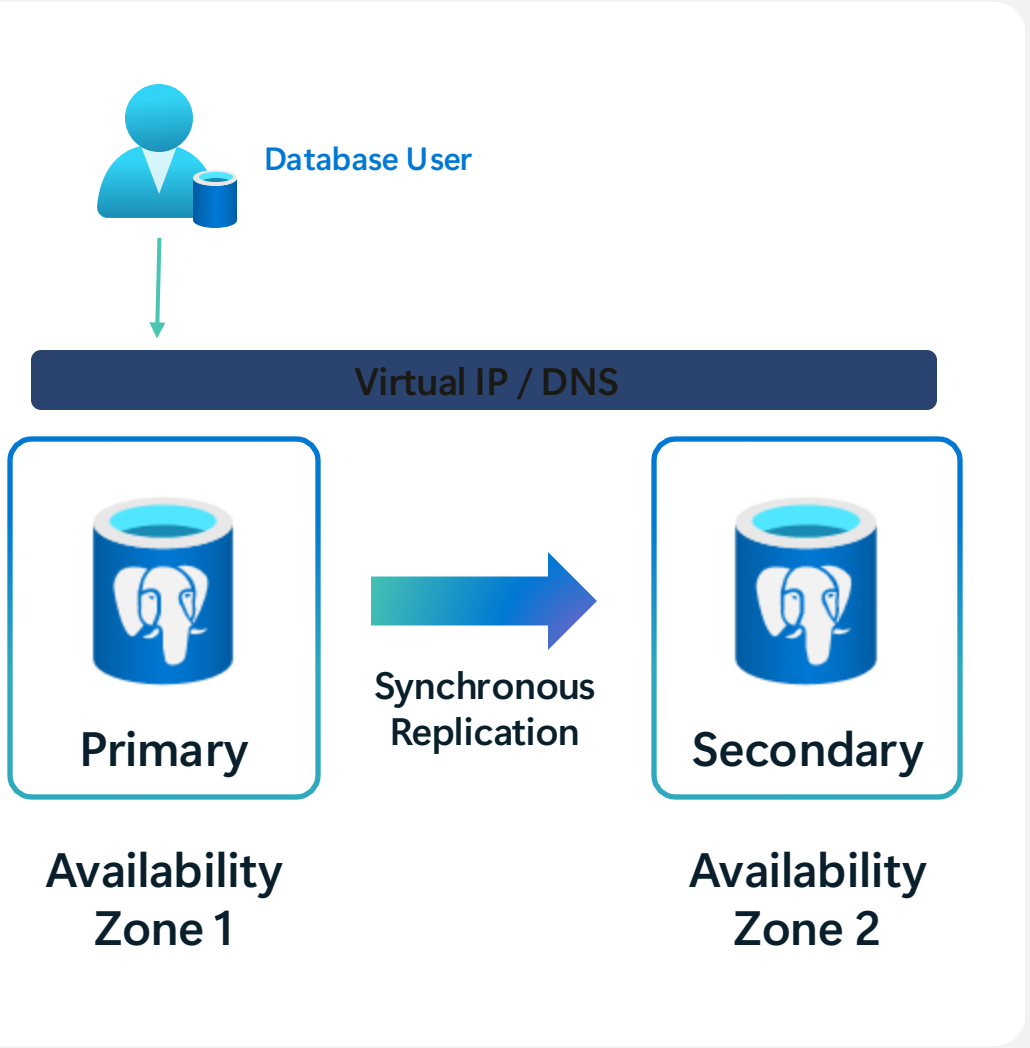


HA-DR demo at Posette Conference 2024:
https://youtu.be/Yeb2Nv5mZtg?list=PLlrxD0HtieHg5ldHlEn3nMMpClX0RUntH&t=549

# Disaster Recovery

**Planned** or **Unplanned** failover to replicas running in any other Azure region

Execute a Planned Failover to **synchronize data between regions** prior to failover

Execute Unplanned Failover for **high impact emergency events** to move workload immediately without data synchronization

Both approached support Virtual IP with no application changes required

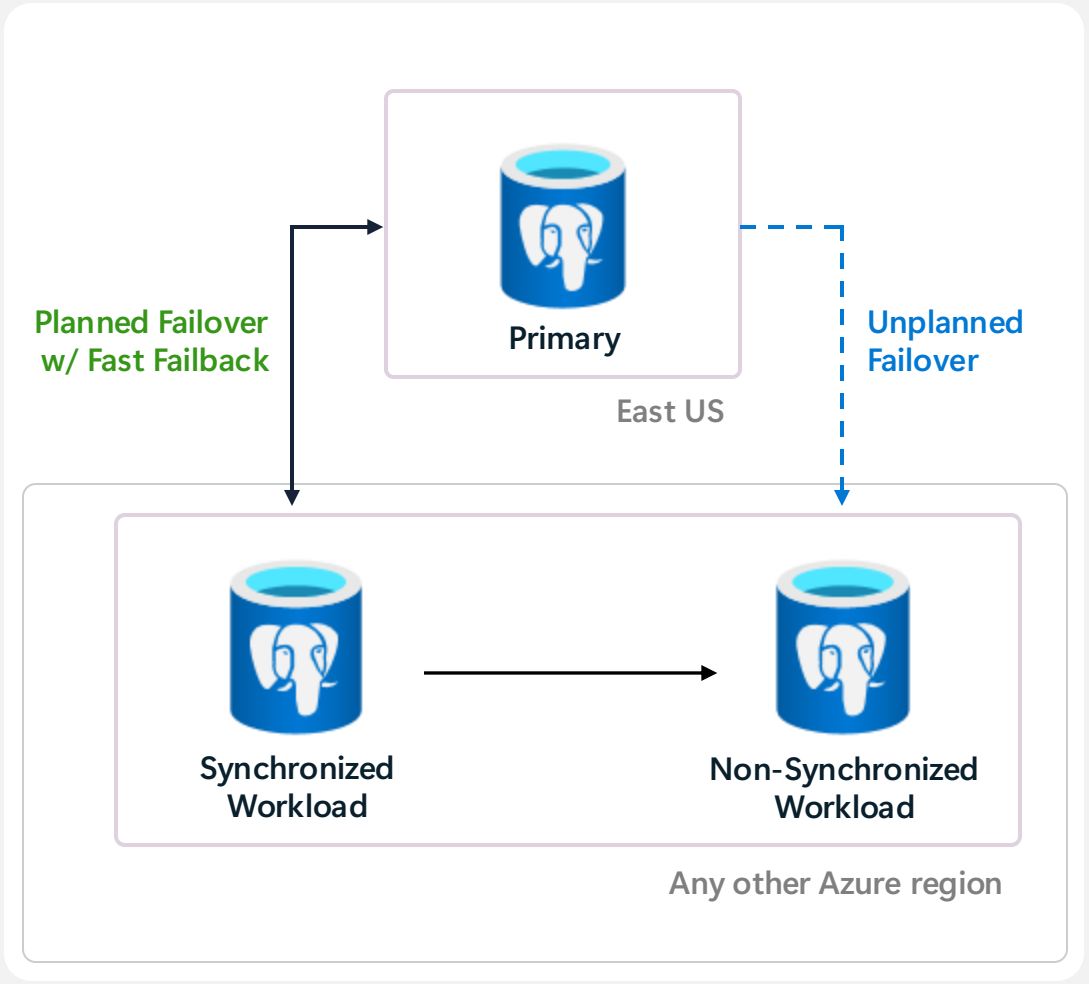

# Disaster Recovery

**Backup protection for major disruptions**

**On-Demand** or **Automated** backups protect data from loss with 30 Day retention

Configure backups for **geo-redundant storage**, and restore to different regions in event of failure

Long-term backup retention for **up to 10 years** supports compliance requirements for data retention

RPO of 5 minutes

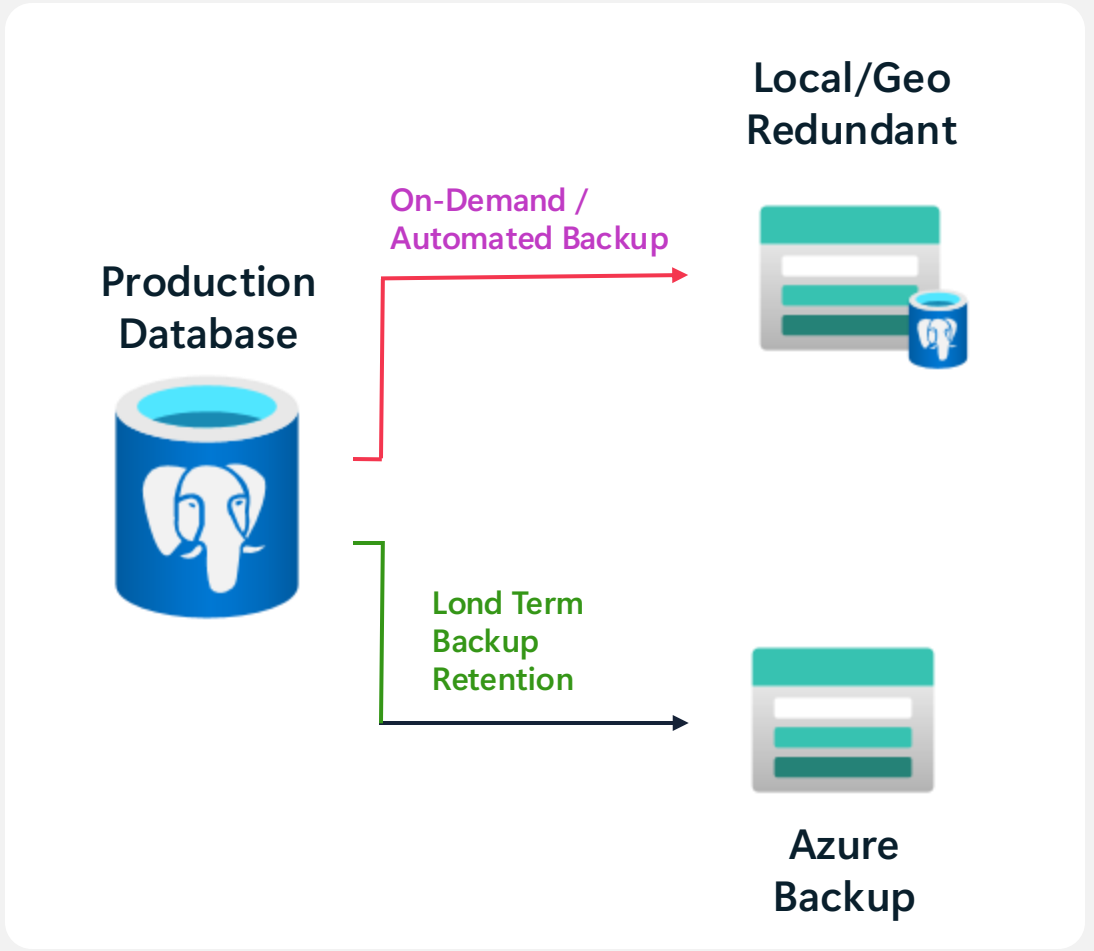

# Comprehensive IaC

**Infrastructure-as-Code**

Use your **preferred IaC provider** to programmatically deploy and manage Flexible Server resources

Integrate with existing **infrastructure workflows**

**Version control** IaC for robust resource change management

## Terraform

Open-source IaC tool for configuring and deploying cloud infrastructure.

## Bicep

A domain-specific language (DSL) that uses declarative syntax to deploy Azure resources

## Ansible

Declaratively manage your Azure using a simple configuration language

## ARM Templates

Declaratively manage your Azure using a simple configuration language

.NET Aspire

Azure-Samples/eShopOnAzure

+ Rest API's, CLI, and PowerShell also

# Built-in Optimization

Unique to Azure

## Index Recommendations

Automatically determine optimal indexes based on user activity

CREATE and DROP recommendations

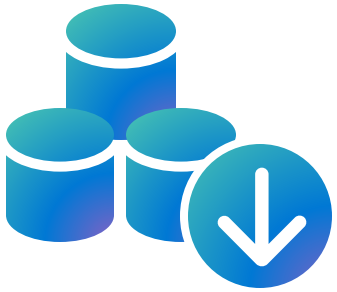

**Minimize Server Resources**

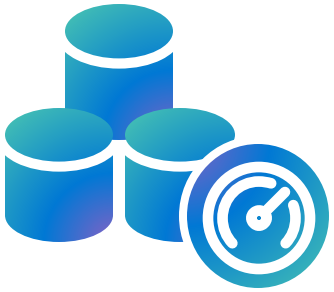

**Maximize Workload Performance**

[Unleash the power of automated index tuning in Azure Database for PostgreSQL Flexible Server - Microsoft Community Hub]

# How Vectors Work: Important in the AI Era

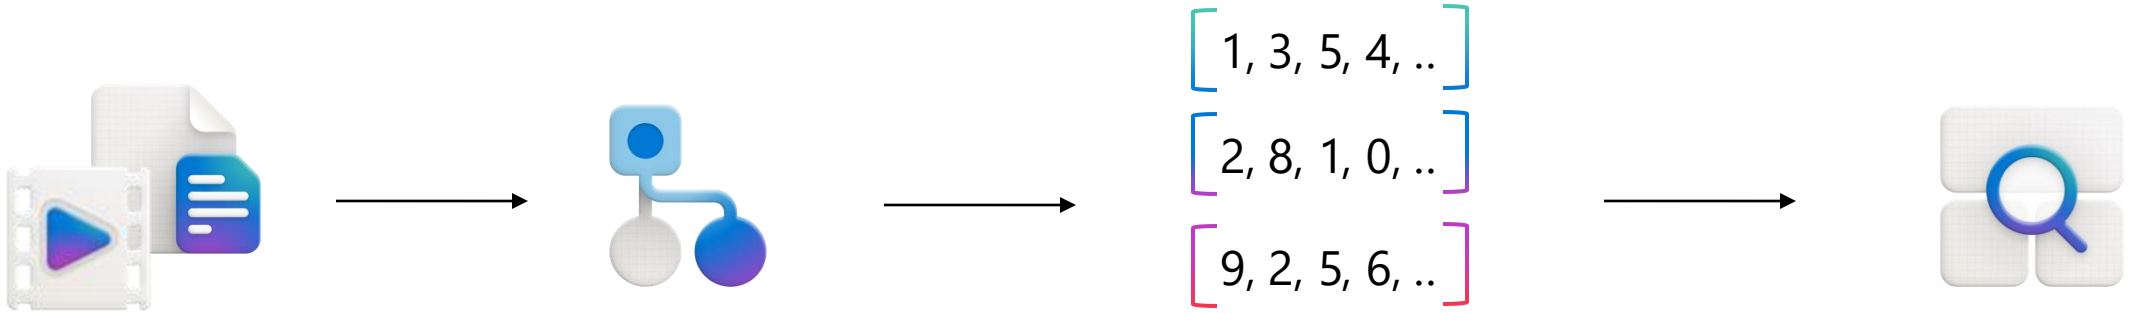


Raw Data → Deep Learning Model → High-dimensional Vectors → Vector Similarity Search

**Vector similarity search empowers Generative AI apps**

Recommendation

Question Answering

RAG apps

...

# Azure Database for PostgreSQL: Native Vector Search

**Open-source Pgvector extension** provides support to store, index, and query vectors for similarity search scenarios

Supports multiple vector **distance functions**

Enables AI solutions to seamlessly integrate into **existing OLTP Postgres apps** without exporting data to specialize systems

Access control, encryption, high availability, disaster recovery all **"just work"**

## Generative AI apps

**RAG (Retrieval Augmented Generation) apps**

Retrieve private data to ground LLM model responses

**Recommendation/Semantic Search**

Retrieve similar documents by distance between vectors

**Hybrid Search**

Combine vector search, row filtering, and full-text search

# Vector indexes supported today

## IVFFlat

- Clusters vectors by applying k-means clustering.
- Memory efficient but requires index rebuilds.

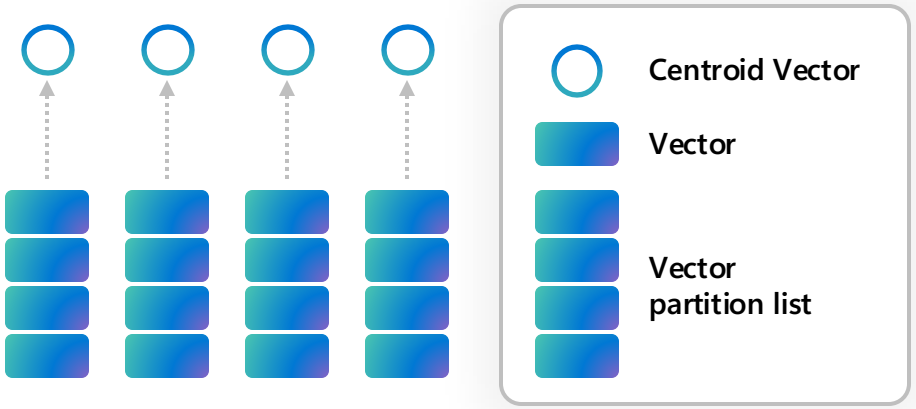


## HNSW

- Builds a multi-layer graph with long and short connections between the vectors.
- The graph can be incrementally updated.

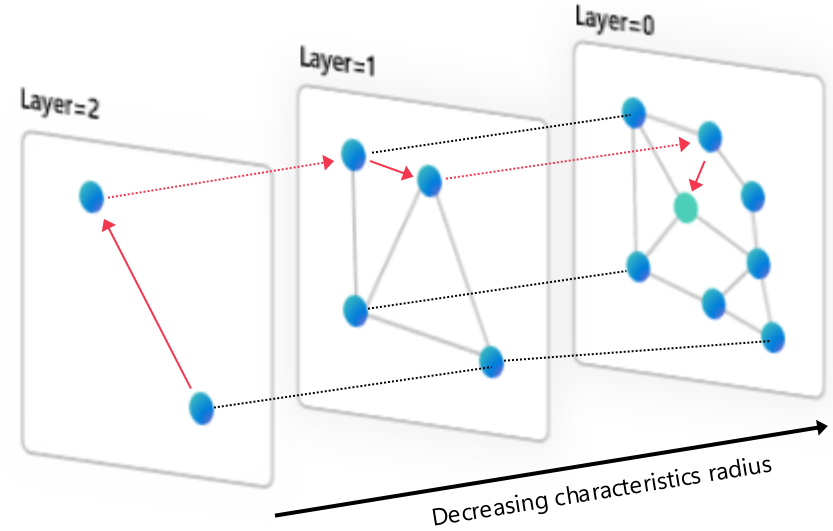

# AI Services integrated with Azure Postgres

**Make remote calls directly from PostgreSQL**

The **azure_ai extension** provides a SQL-based interface to integrate with AI services

Supports:

- Azure OpenAI
- Azure AI Language Services
- Azure AI Translator
- Azure Machine Learning

Enables developers to **rapidly integrate AI capabilities** into their app without complex re-architecture or refactoring

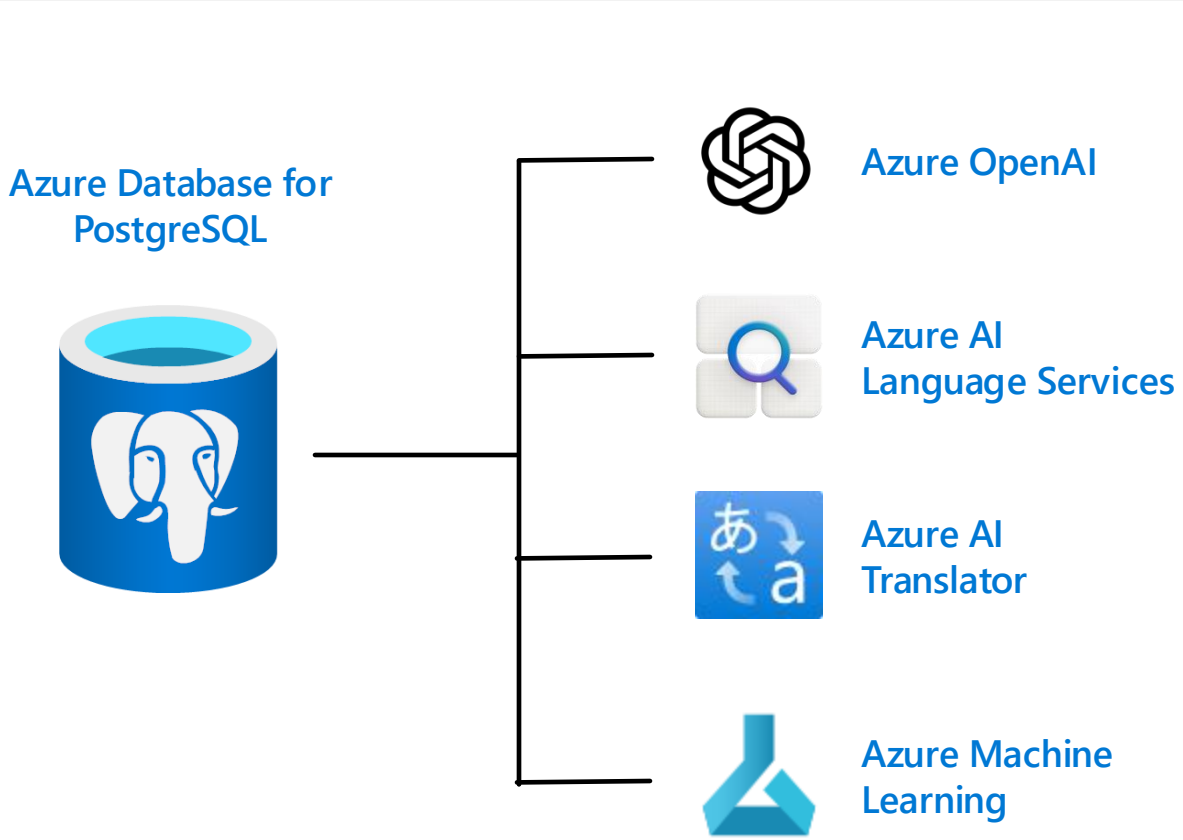


[Azure AI extension on Azure Database for PostgreSQL - Flexible Server]

# In-Database Embedding Models

## Low-latency embedding creation for OLTP workloads

The **azure_local_ai extension** enables vector embeddings to be generated locally within the Postgres server

Based on the Microsoft **open-source E5** embedding model

Benefits:

- **~10x faster** creation time
- No external service setup, maintenance, or transaction costs
- All data remains **within Postgres**
- Perfect for workloads where the **underlying data changes frequently**

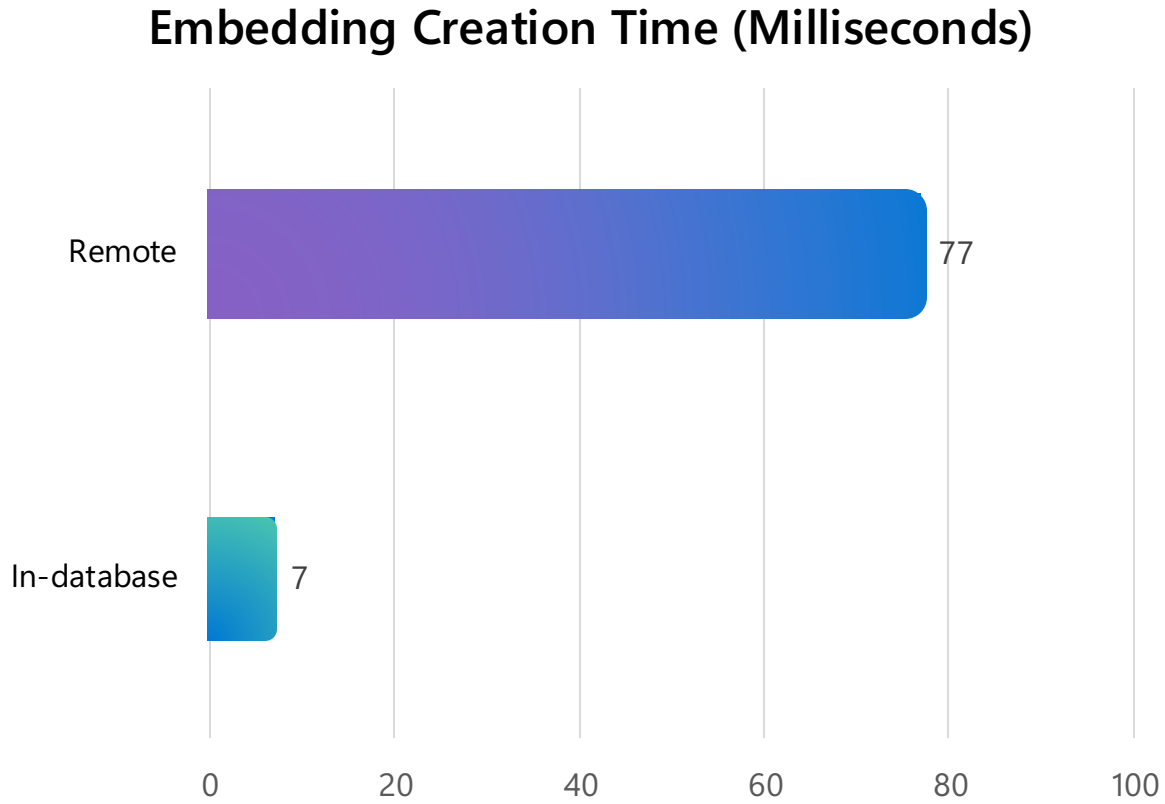


[Create in-database embeddings with azure_local_ai extension - Azure Database for PostgreSQL - Flexible Server | Microsoft Learn]

# Vector Generation

## Unique Remote + In-Database Embedding Models

**Remote Embedding Models**

```
SELECT * FROM <table>
ORDER BY
database_description <->
azure_openai.create_embeddings(
'text-embedding-ada-002',
'Databases with vector support')
```

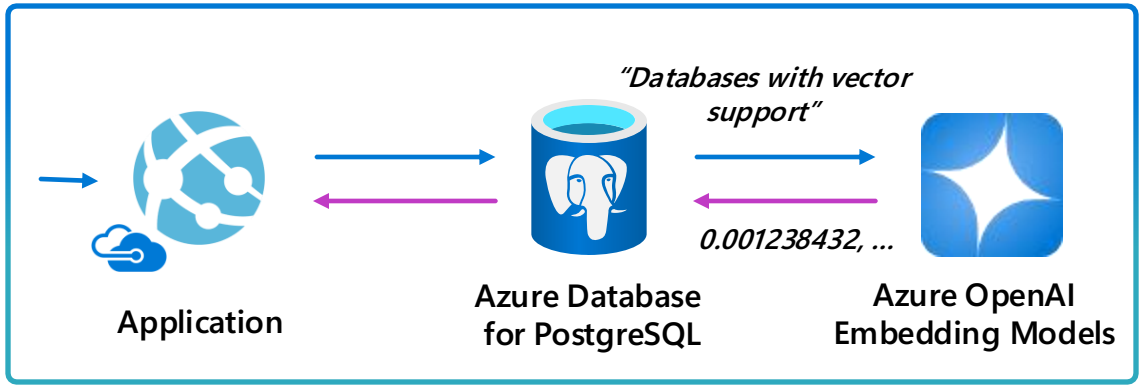


**In-Database Embedding Models** (Preview)

```
SELECT * FROM <table>
ORDER BY
recipe_embedding <#>
azure_local_ai.create_embeddings(
'multilingual-e5-small:v1',
'Databases with vector support')
```

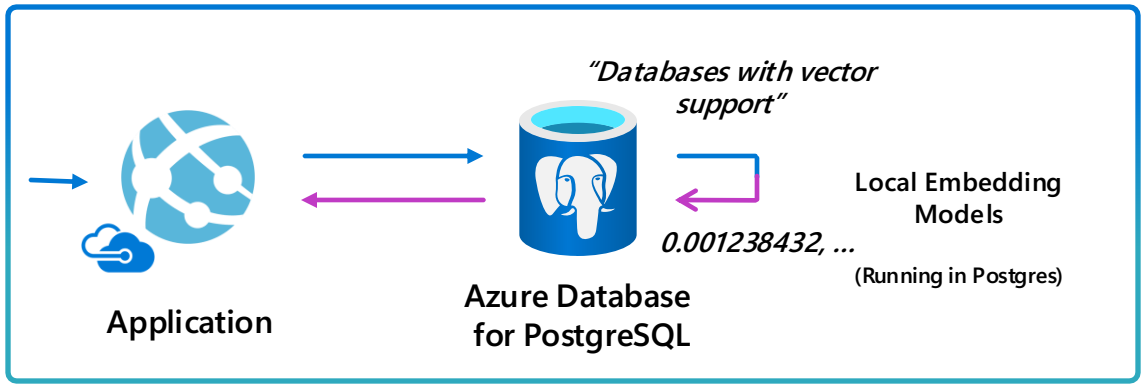

## Announcing

# Mirroring Azure Database for PostgreSQL Flexible Server in Fabric!

## Private Preview

Private preview form

https://forms.office.com/r/H0Mq1Yz8Pw

# Mirroring Azure Database for PostgreSQL Flexible Server

## Supercharge analytics time to value

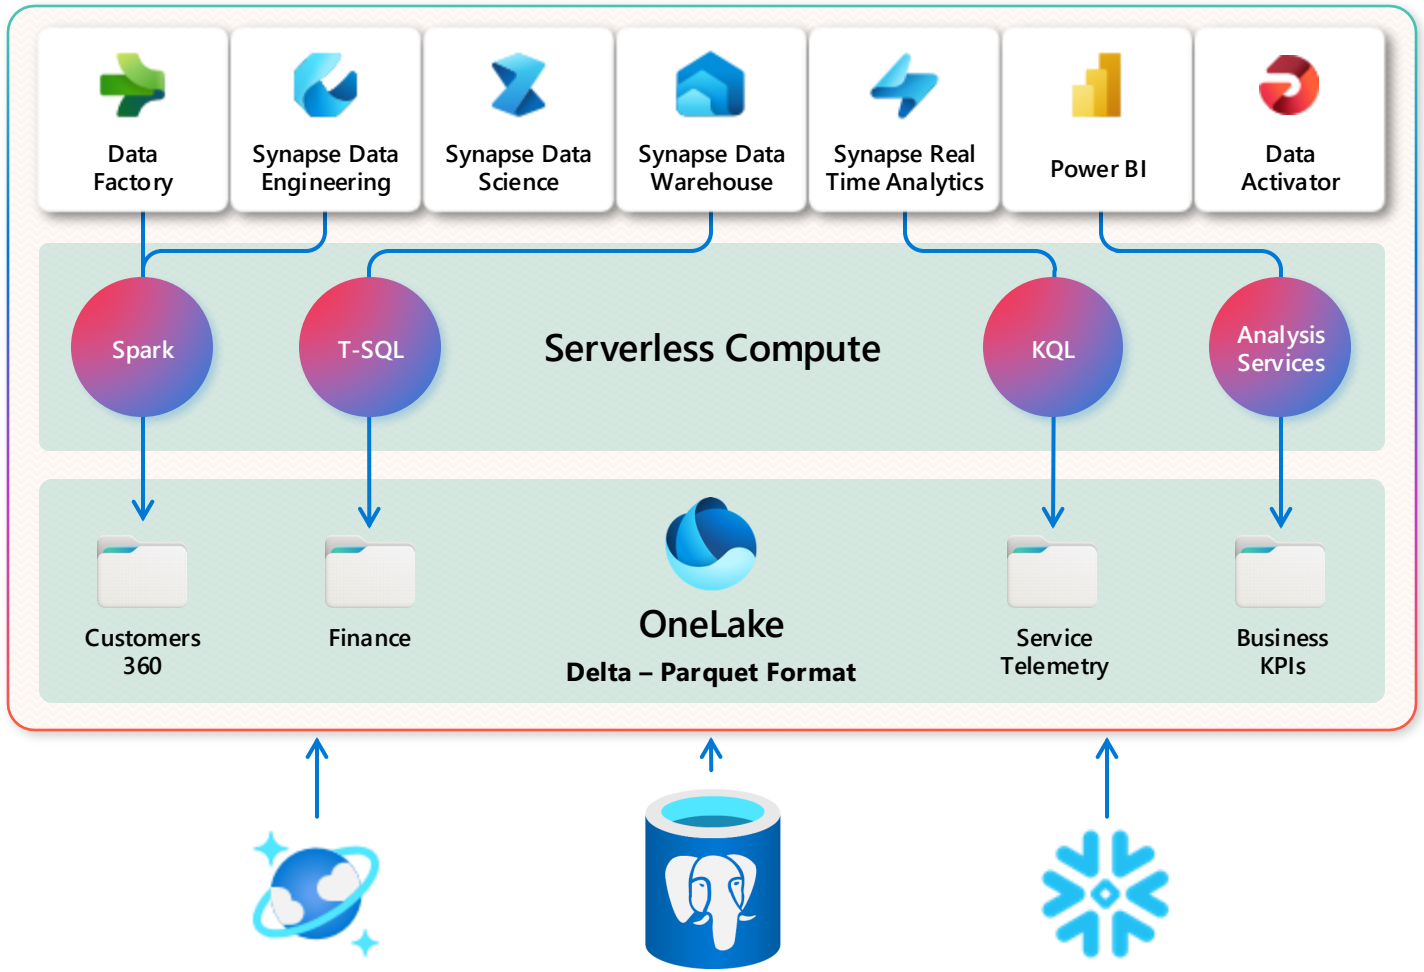


Mirroring replicates databases to Fabric with **zero ETL**

Compute replication is included with your Fabric capacity for **zero cost**.

**Free storage** costs based on Fabric capacity

Data is replicated into One Lake and **kept up-to-date in near real-time**

Mirroring **protects operational databases** from **analytical queries**

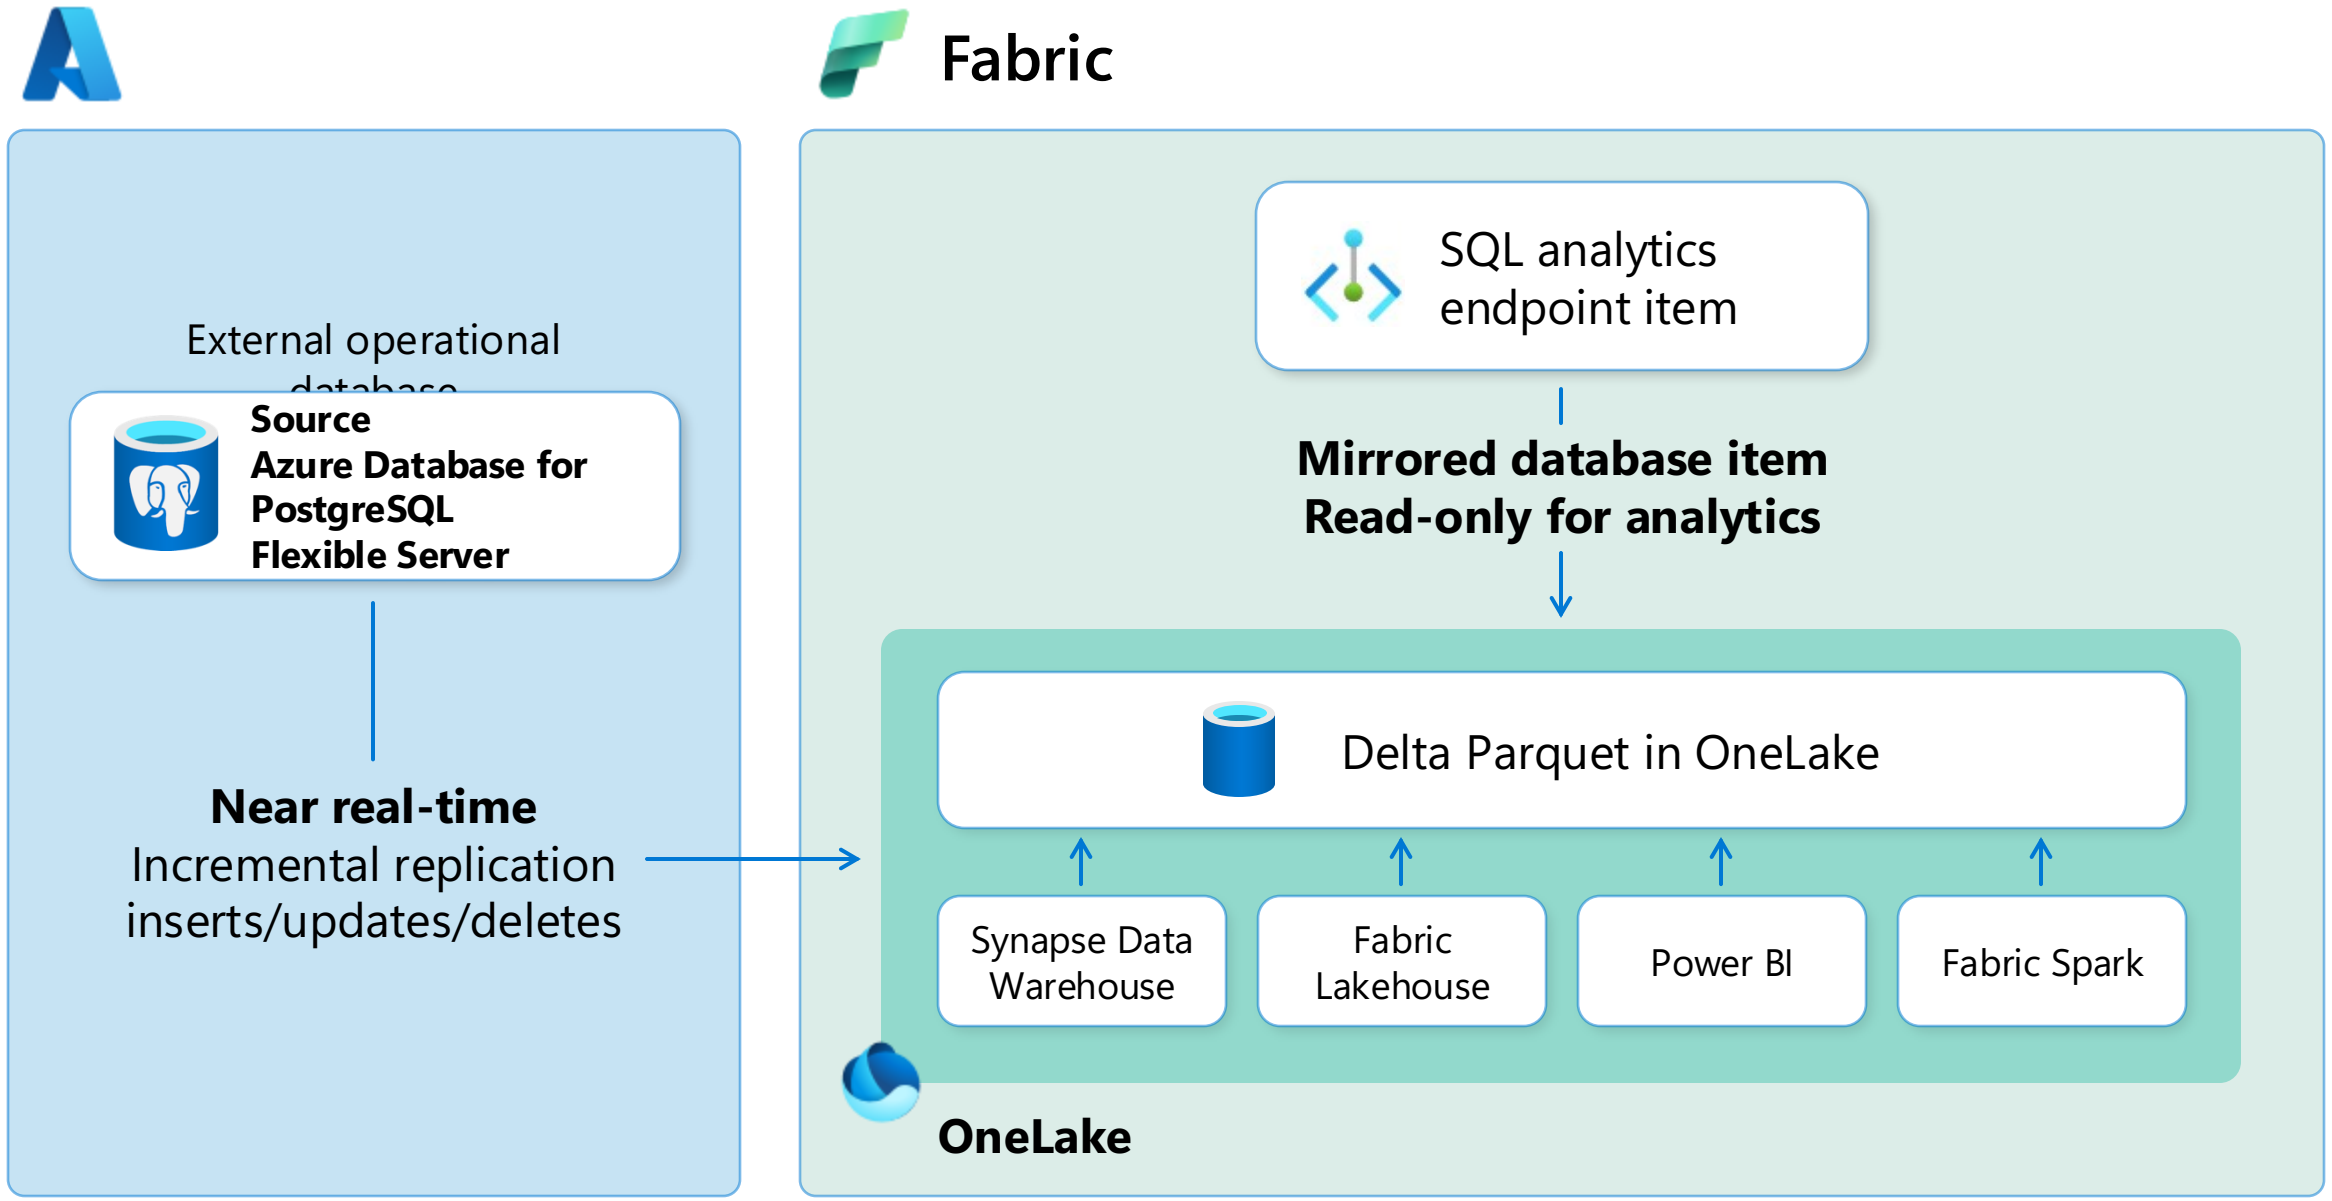
Fabric
External operational database
Source
Azure Database for PostgreSQL Flexible Server
Near real-time
Incremental replication
inserts/updates/deletes
SQL analytics endpoint item
Mirrored database item
Read-only for analytics
Delta Parquet in OneLake
Synapse Data Warehouse
Fabric Lakehouse
Power BI
Fabric Spark
OneLake

# Fabric Experiences Unlocked with Mirroring

## Cross-database querying with multi-cloud Shortcuts

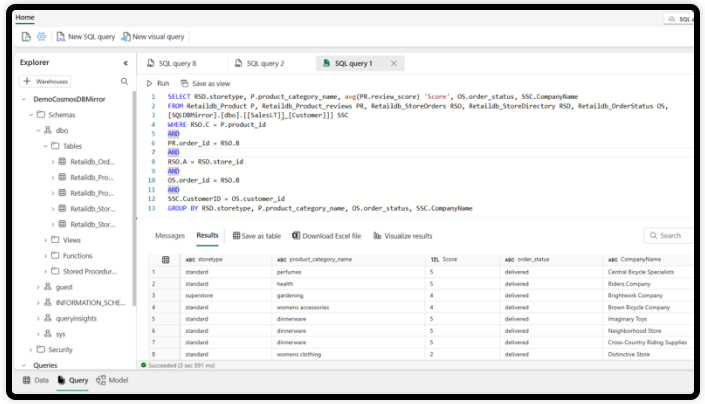


## Data Science, ML & AI experiences

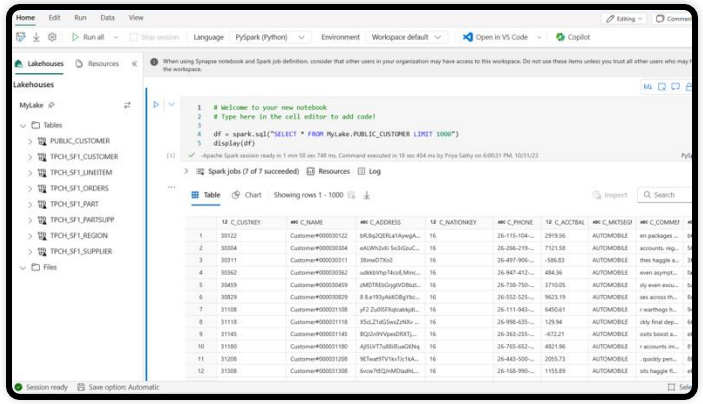


## Power BI Direct Lake mode

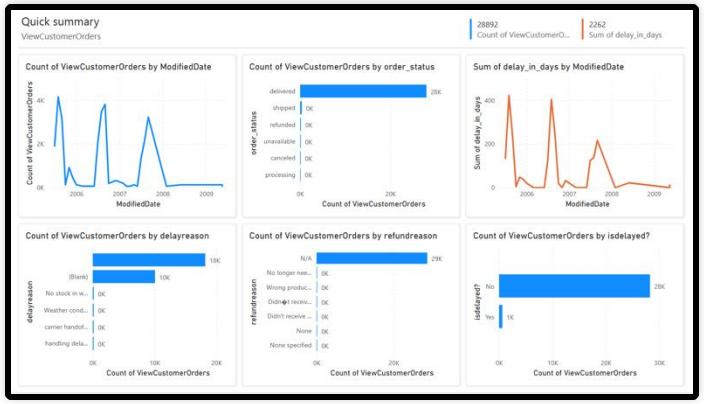

# PostgreSQL Database Change Data Capture (CDC) source connector for Microsoft Fabric

## Public Preview

Add PostgreSQL Database CDC source to an eventstream - Microsoft Fabric | Microsoft Learn

# Real-Time Intelligence in Microsoft Fabric

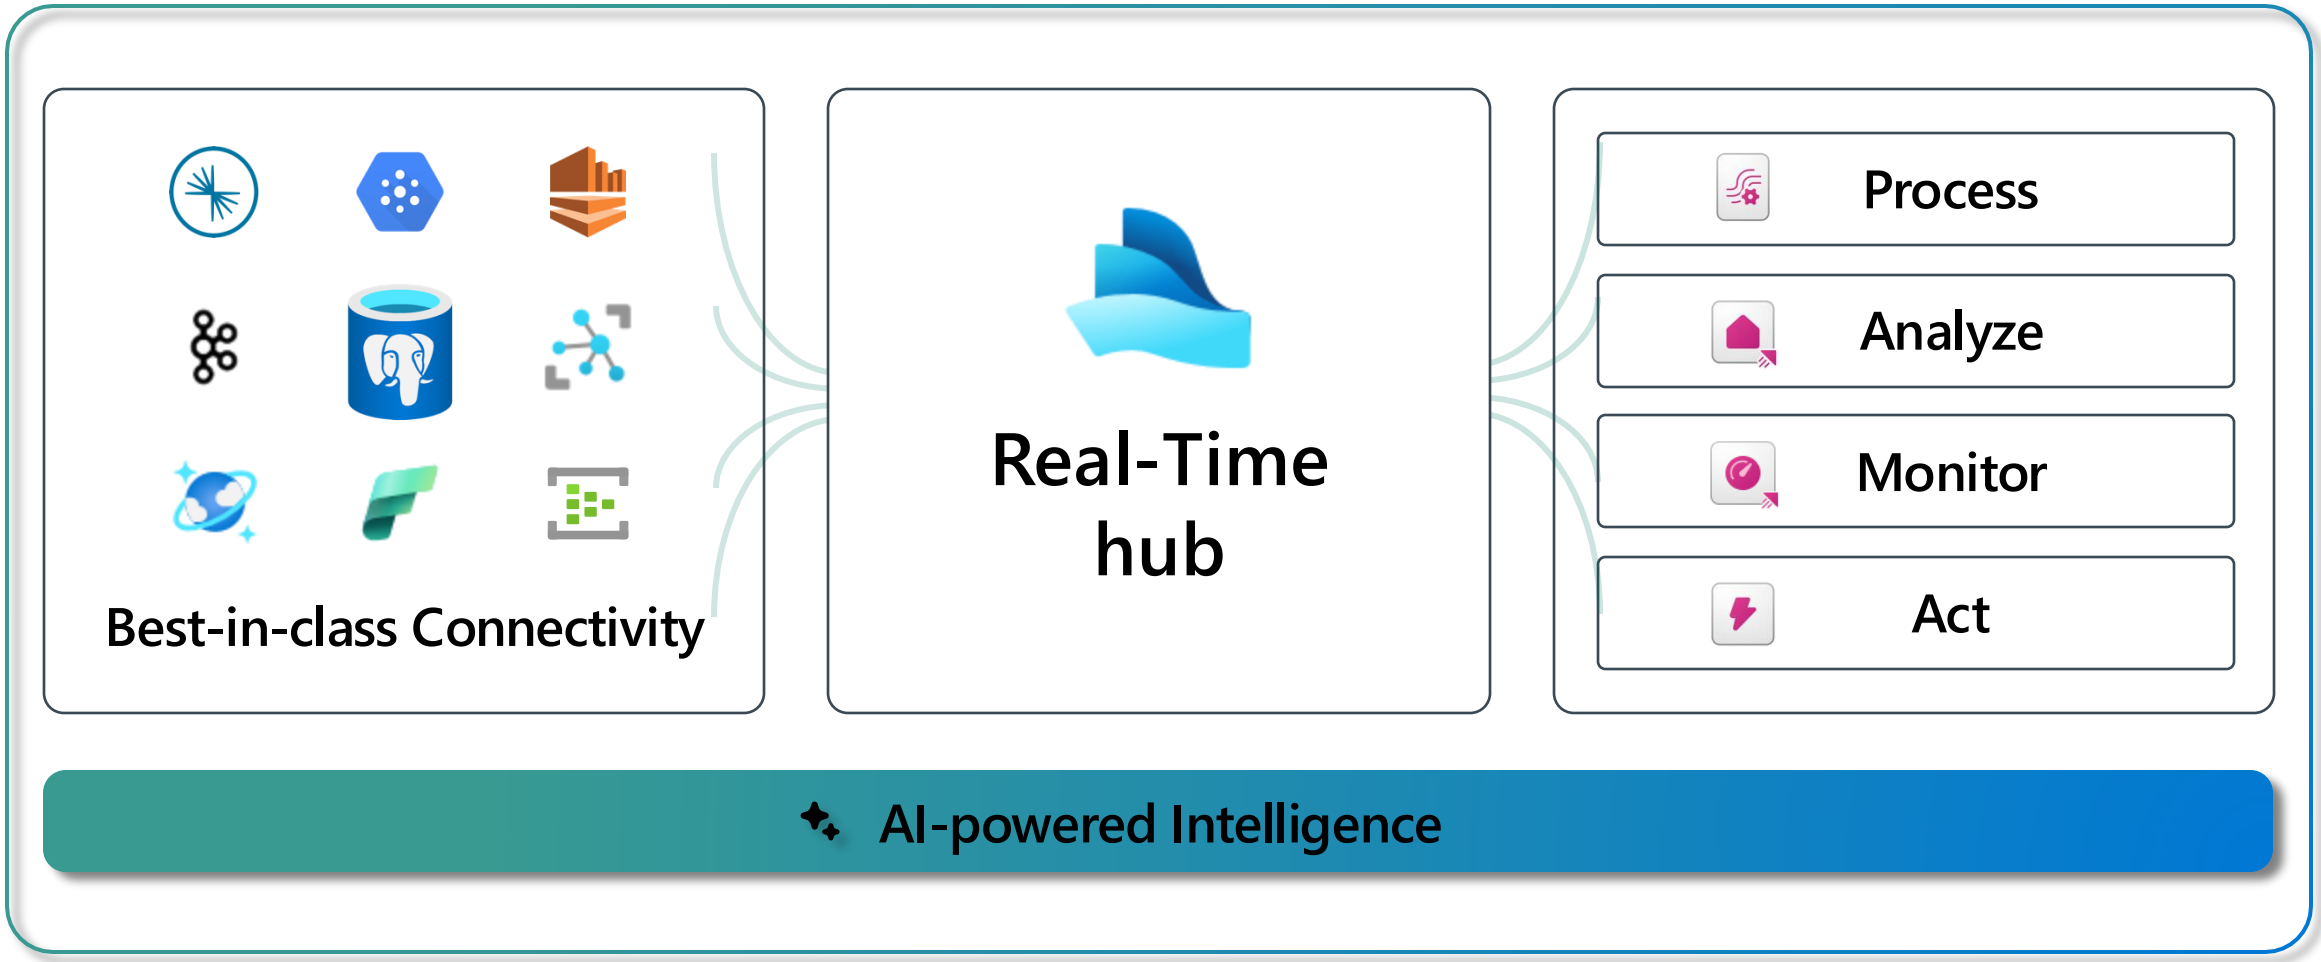

# Real-Time Intelligence with PostgreSQL data

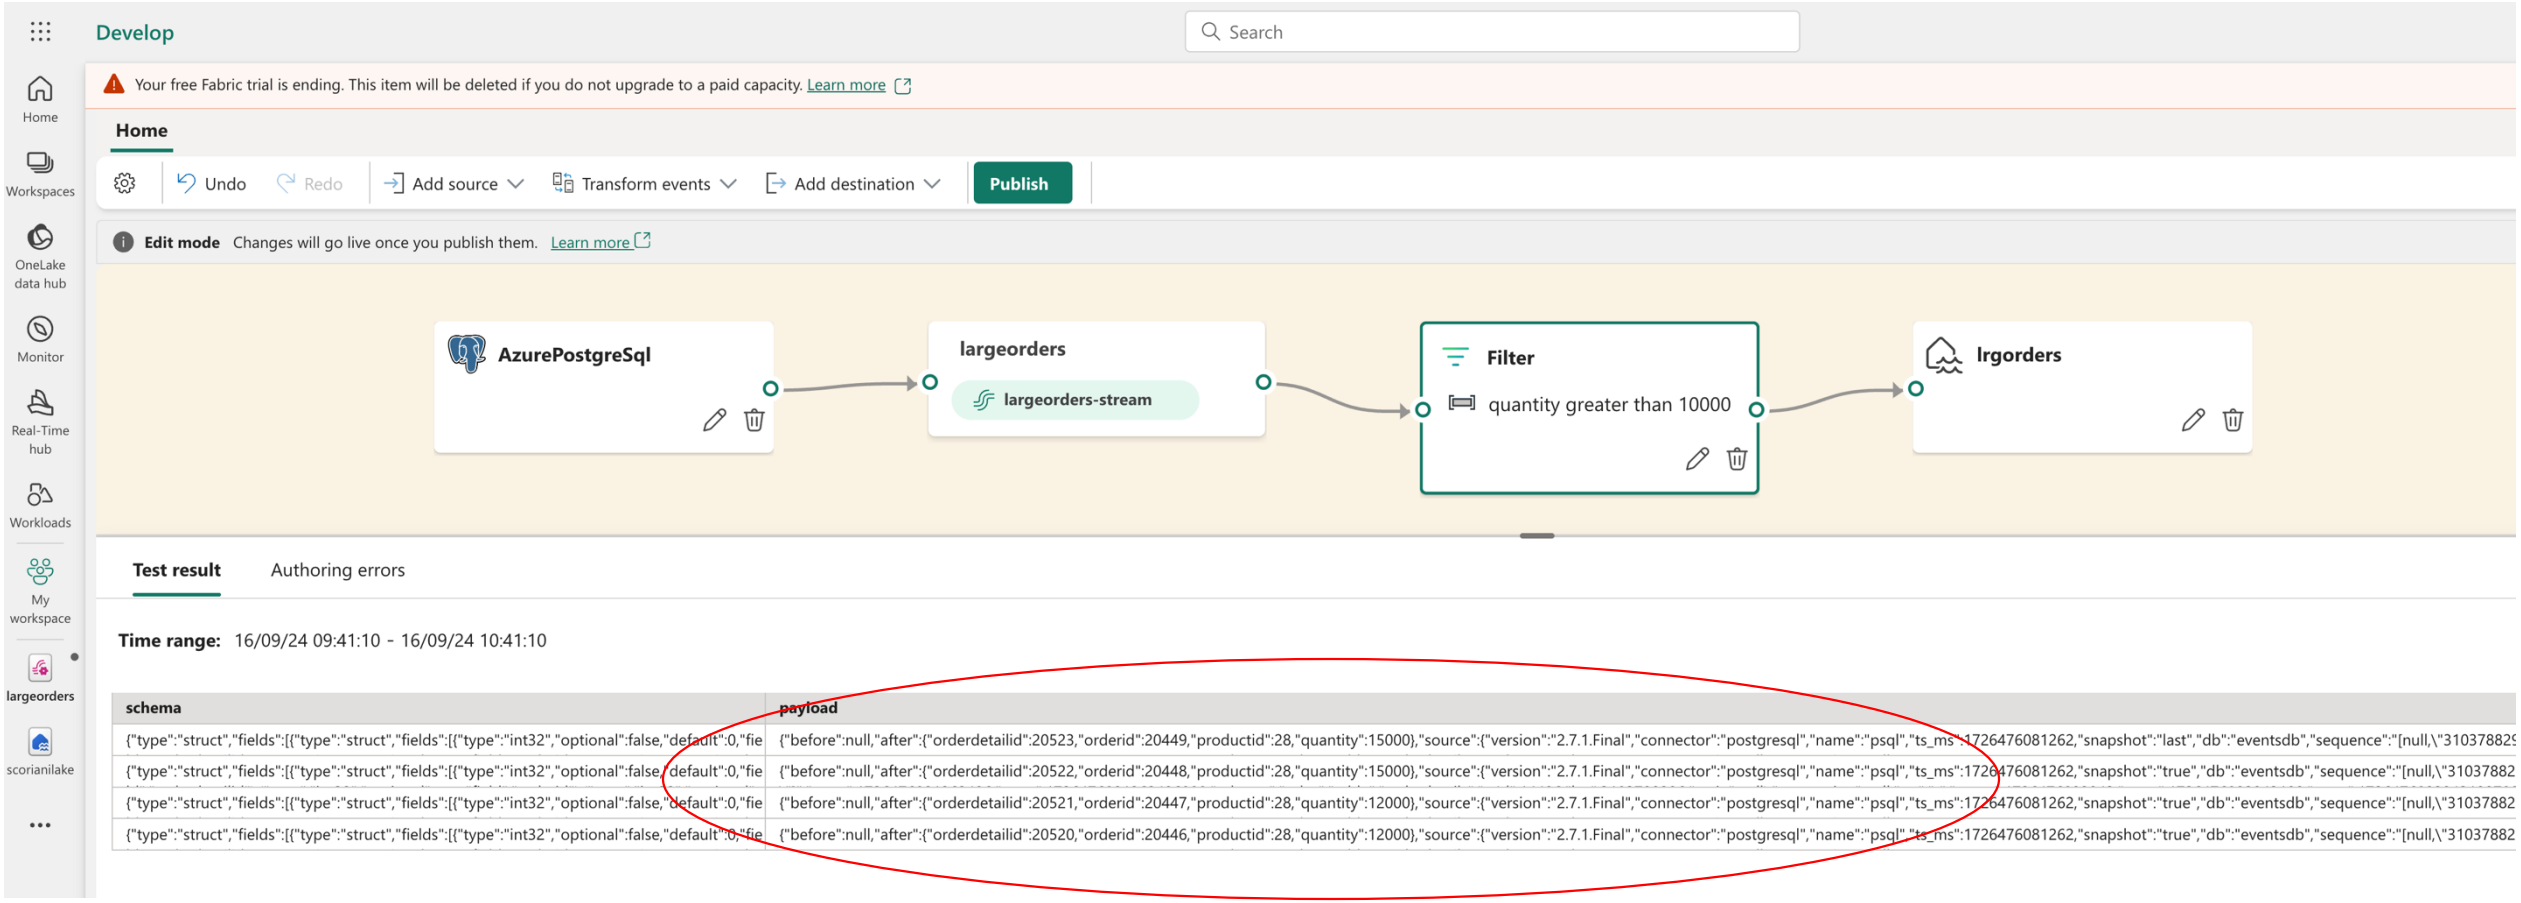

# Built-in migration tools

## Migrate from Postgres anywhere to Flexible Server with ease

PostgreSQL – Single Server

Azure VM

On-Premises

Amazon EC2

AWS RDS

### Migration service available for easy PostgreSQL migration

- Built-in as a managed service
- Secure and Reliable
- Optimized for performance
- Initiate migration via Azure Portal or CLI
- Connection established from target to source
- Performs offline/online database migration

### Azure Database for PostgreSQL – Flexible Server

- Full Azure AI Integration
- Automatic scalability
- Intelligent Performance tuning
- Automatic High availability
- No operational overhead
- Automatic Disaster Recovery

*More sources coming soon (GCP & AWS Aurora coming in few weeks)

# Azure PostgreSQL Resources

**Azure Database for Postgres homepage**
http://aka.ms/postgres

**Azure Database for Postgres Docs**
http://aka.ms/postgresdocs

**Azure Postgres on X**
@AzureDBPostgres

**Azure Postgres Feedback Forum**
https://aka.ms/pgfeedback

**Get Started for Free with an Azure Free Account**
https://aka.ms/try-postgres-free

**Azure Database for Postgres Blog**
https://aka.ms/azurepostgresblog

**Azure Postgres on LinkedIn**
https://www.linkedin.com/company/azure-database-for-postgresql/

**AskAzurePostgreSQL**
AskAzurePostgreSQL@microsoft.com

Microsoft

# Thank you